

STAT 652: Predicting Flight Delays Project

Vincent Chiu

11/26/2019

Contents

| | | |
|-----------|---|-----------|
| 1 | Introduction | 1 |
| 2 | Data | 2 |
| 3 | Methods: | 2 |
| 3.1 | Data Preprocessing | 2 |
| 3.2 | Exploratory Analysis | 3 |
| 3.3 | Principal Component Analysis (PCA) | 3 |
| 3.4 | Cross Validation | 3 |
| 3.5 | Models | 3 |
| 4 | Results | 4 |
| 4.1 | GBM | 4 |
| 4.2 | Linear Regression | 5 |
| 4.3 | Comparison | 5 |
| 5 | Conclusion and Discussion | 5 |
| 5.1 | Discussion | 5 |
| 5.2 | Conclusion | 5 |
| 5.3 | Future Work | 5 |
| 6 | Code | 6 |
| 6.1 | Preparing the programming environment | 6 |
| 6.2 | Data Preprocessing | 6 |
| 6.3 | try features scaling | 11 |
| 6.4 | Exploratory Data Analysis | 14 |
| 7 | predicting 0 | 18 |
| 8 | predicting the mean | 19 |
| 9 | predicting the median | 19 |
| 10 | linear regression with dep | 19 |
| 11 | gbm | 20 |
| | References | 28 |

1 Introduction

The goal of this project is to predict the response variable, departure delays for a particular flight given the explanatory variables.

2 Data

The dataset consists of information about all the flights leaving from New York City in 2013. The dataset contains 43 variables in total. The dataset is an amalgamation of several datasets including datasets containing information on weather, the airports, the flights, and the models of airplanes. The training dataset provided to us contains 200,000 observations. Please see table 1 for what the data looks like.

3 Methods:

We will now outline the various methods used to clean and perform prediction on the data. We will discuss our techniques for data preprocessing, and cross validation, and the different models that we tried.

3.1 Data Preprocessing

I performed data preprocessing on the nycflights13 dataset. My data preprocessing steps include the following:

- loading the data from a csv
- setting the random seed for reproducibility of results
- casting all the columns with the character data type into the factor data type
- converting the shed_arr_time and sched_dep_time columns into the POSIX time format so that I can accurately take the difference of them.
- Dropping columns that contain data from after the planes' departure which may leak information about the response variable dep_delay. We drop columns "dep_time", "arr_time", "air_time", and "arr_delay".
- We drop column "year.x" because all the values are 2013
- We also drop tailnum because it produces too many dummy variable columns for one hot encoding.
- Dropping columns which consists of over 50% NAs which include the speed column. However, it should be noted that a rule of thumb suggested by Professor McNeney is to drop any columns with over 5% NAs. We use a different threshold for dropping columns leading to us keeping columns such as model instead of dropping it.
- Afterwards, I impute the missing values. However, there are limitations to this approach of imputing the missing values. It is possible, that the missingness of the plane model variable is related to dep_delay. In this scenario, we may be creating an inferior feature set by keeping the variable 'model' and imputing it. For example, say a highly unreliable plane model that frequently causes long delays has a high probability of being labeled as NA and represents the majority of NAs in the dataset. We would not be able to capture the relationship between this plane model and dep_delay if we imputed the model with the mode. To counteract this affect, we imputed NAs for the remaining columns using the imputeMissings library, adding a Boolean flag which indicates 1 if the associated value was 1 and 0 otherwise. For example, the "model_flag" for a given row is 1 if the "model" value was NA for that given row. Hence, no information is lost from our imputation.
- Normalizing the data to work well with methods like lasso regression.
- Only kept data which had a departure delay of less than 30 minutes late, which reduced the dataset from 200,000 rows to approximately 170,000. This is because we consider extreme delays of over 30 minutes late to be freak accidents which cannot be accurately predicted by the available explanatory variables.

3.2 Exploratory Analysis

3.2.1 Correlations

First, we created a correlation plot for the numeric variables to see if there any correlations between the variables.

We see that there is very little correlation between the response variable `dep_delay` and any of the other variables. Some of the strongest correlations include the correlation between distance and longitude and time zone and a smaller correlation between distance and latitude. This makes sense as most of the planes are inter US flights from west to east or vice versa, there is not as much distance flown in the north south direction. Please see Figure 1 for the correlation plot (Click on the number after Figure to jump to plot).

3.3 Principal Component Analysis (PCA)

Next we performed PCA on only the numeric variables as techniques to perform PCA on mixed datasets (numerical and categorical) was not covered in class. When looking at the contribution of each variable to the first principal component, we notice that the variables `lon`, `distance`, `tz`, `seats`, `alt`, `sched_air_time` have the greatest absolute coefficients for the first principal component. The fact that the aforementioned variables have large coefficients in the first principal component suggests that they are highly correlated with each other. The fact that `dep_delay` has a small coefficient in the first principal component suggests that `dep_delay` is not highly correlated with any of the above variables.

As expected, it turns out that variables like `lon`, `distance` and `tz` are not important for predicting `dep_delay` according to the `gbm` model. This maybe be because although variables like `lon`, `distance` and `tz` help explain most of the variance in the dataset, they have a weak relationship with `dep_delay`.

Please see table 2 for the proportion of variance explained by each principal component. Please see table 3 for the coefficient of each variable for each principal component ordered by magnitude of coefficient.

3.4 Cross Validation

Initially, I used the most basic cross validation technique where I have a training dataset and a cross validation dataset. I split the original data into a ratio of 2/3 train and 1/3 of the data for cross validation. There is a additional data which would be provided by the professor at a later date which we will use as the holdout test set. I believe that 2/3 of the data gives enough data for the models to train on while 1/3 is enough data for us to get an accurate assessment of the error. k-folds cross validation was not initially used in order to save on compute time as we were initially only exploring the models. k-folds cross validation would increase training time for the models by a factor of k. However, k-folds cross validation would lead to a more stable estimate of holdout test set error.

3.5 Models

We first explored some basic models to establish a baseline performance and compared it to our most sophisticated model, the Generalized Boosted Regression Model (GBM).

3.5.1 Basic Models

`dep_delay` is the number of minutes that the plane either departs early or late. Negative numbers are for early departures and positive numbers are for the number of minutes the plane is late. First, I used a basic model of simply predicting the `dep_delay` to always be 0. This was done to establish baseline performance.

This model had an root mean squared error (RMSE) of 8.30571. TODO The model in which I predicted the mean for all the predictions had an RMSE of TODO.

3.5.2 Linear Regression

A linear regression model assumes a linear relationship between the explanatory variables and the response variable. The model minimizes the squared loss function. This minimization process generates coefficients which are used for a linear combination of the explanatory variables. The linear combination is the prediction. Then I tried linear regression with `dep_delay` as the response variables and all the other remaining variables as the explanatory variables. This model was better than predicting the mean with an RMSE of TODO. This suggests that there is some relationship between the `dep_delay` and the explanatory variables.

3.5.3 Generalized Boosted Regression Model (GBM)

How boosted regression models work is we first start off with a baseline prediction. For example, we can use the mean as our baseline prediction. Then we use a base classifier to iteratively predict on the residuals multiplied by the shrinkage hyperparameter. Then this base classifier is added to the ensemble of base classifiers trained so far. A lower shrinkage effectively means a lower learning rate and therefore you need more iterations to reduce the train set residuals by the same amount. The benefit of a smaller shrinkage (with sufficient trees) is that you end up with a larger ensemble of trees that can reach a lower cross validation loss. This iterative prediction process is called boosting. In our case, our base classifiers are regressor trees. Each tree decides on its splitting criterion greedily by picking the split which results in the lowest mean squared error or some other splitting heuristic. We continue for n number of trees where n is specified by the user. Each iteration produces one tree, so the number of iterations is equal to the number of trees. Afterwards, we tried a Generalized Boosted Regression Model (GBM). This model had the lowest RMSE of 7.94458 on the cross validation set after it was tuned to have a shrinkage of 0.01 and around 16,000 trees. Shrinkage is proportional to the learning rate. 16,000 trees is the number of trees used in the model. Each iteration uses 1 tree, so 16,000 trees also refers to the number of iterations. According to the vignette, the RMSE can always be improved by decreasing shrinkage, but this provides diminishing returns. A good strategy would be to pick a small shrinkage that balances performance and compute time. Then with this fixed shrinkage value, increase the number of trees until you get diminishing returns. We decided to follow the aforementioned strategy. Please see table @4 for a table of the relative influence of each explanatory variable. Here, you can see the relative influence for each variable for gbm.

For a gbm, the improvement in the splitting criterion (which is mean squared error for regression) for a given variable is calculated at each step. The relative influence for a given variable is the average of these improvements over all the trees where the aforementioned variable is used.

4 Results

In regression and gbm, I found different features to be important.

Please see table @5 for all the models and their root mean squared error (RMSE) on the cross validation set.

4.1 GBM

For the best gbm model, `dest` which refers to which airport a given plane was flying to was the most important feature. However, the one hot encoding versions of carrier were the most important features for regression. Based on the relative influence scores provided by the gbm, some of the most important feature variables include `dest`, `model`, and `sched_dep_time_num_minute`. The `dest` column contains the airport code for where a given flight is flying to. Based on my run of gbm with a shrinkage of 0.01 and 16834 trees, `dest` was the most important feature with 49.56 relative influence. (“Gradient Boosting Machines · UC Business Analytics

R Programming Guide” 2019). `dest` is the destination airport code. `'sched_dep_time_num_minute'` is the number of minutes since the beginning of a given day for that flight. `'model'` is the plane model.

4.2 Linear Regression

On the other hand, `dest` does appear as an important feature in linear regression as well but it is not the most important feature. I surmise that if we can somehow sum up all the contributions from each of the one-hot-encoded variables derived from `dest` then, it might appear as the most important feature for linear regression as well. We can try using ANOVA in order to measure the statistical significance of `dest`. Performing ANOVA on comparing linear regression model with and without `dest`, it was determined that due to the low p-value of 0.0001863 associated with having `dest` that keeping at least one of the one hot categorical variables derived from `dest` is beneficial for the linear regression model.

4.3 Comparison

`Dest` was most important feature in `gbm`. It is possible that `dest` is important in combination with other variables which is something that the linear model without interaction terms cannot capture the relationship of whereas `gbm` can discover these non linear relationships.

TODO: try interaction terms , try anova.

5 Conclusion and Discussion

5.1 Discussion

We considered removing outliers in terms of `'dep_delay'` in train but not in test, then use k-folds cross validation on test to determine how many outliers we should remove to boost performance on the cross-validation set. We considered removing highly influential points in order to train a better model. In this case, we consider highly influential points to be points with high cook's distances. However, this was infeasible as we did not have enough computational resources available and it took too long.

5.2 Conclusion

None of the models that we tried performed particularly well. We surmise that this may be due to the explanatory variables having a weak relationship with the `'dep_delay'` variable. There is a lack of information about a particular flight before it reaches NYC. Instead, we get information about where the flight is going next which through commons sense, would reveal less information about the current condition of the plane and what kind of maintenance it would need and therefore what `dep_delay` it would have. Out of the methods that we covered in class, I found gradient boosted models to provide the best performance based on having the lowest root mean squared error on the cross validations et. I believe that this makes sense because GBMs are able to capture non linear relationships between the explanatory variables and `'dep_delay'` whereas linear regression cannot.

5.3 Future Work

TODO: add `dep_delay` vs explanatory plots. TODO: remove points that are outliers ie `dep_delay > 200` or `300` etc. or remove less than x number of points. then use k-folds cross validation on cross validation set where no points were removed. can repeat k-folds for different seeds. can just try this on my quickest model, i.e. linear regression. should be bowl shape vs RMSE vs. number of points removed. theoretically

I also considered removing based on cook's distance but this took too long to compute.

5 folds with 10 different random seeds

have train, CV and test set 1/3 train, 1/3 CV, 1/3 test 2/3% train, 1/3%CV, wait for prof test set

try lasso regression add more insights into generating dataset.

Conclusion stuff. certain planes tend to fly back and forth and that is why dest is good? interaction terms may be important because they are hard to predict in advance when you schedule flights which leads to delays? add more about shortcomings of approach. cut off too much of tail, reduce predictive power on test dataset which has tail in exchange for better performance on stuff not in tail.

maybe for stretch goal try rank stuff that prof did.

6 Code

6.1 Preparing the programming environment

6.1.1 Loading Libraries

```
library(tidyverse)

## -- Attaching packages ----- tidyverse 1.2.1 --
## v ggplot2 3.2.1      v purrr  0.3.3
## v tibble  2.1.3      v dplyr  0.8.3
## v tidyr   1.0.0      v stringr 1.4.0
## v readr   1.3.1      v forcats 0.4.0

## -- Conflicts ----- tidyverse_conflicts() --
## x dplyr::filter() masks stats::filter()
## x dplyr::lag()     masks stats::lag()
```

6.2 Data Preprocessing

6.2.1 Loading the data

```
library(nycflights13)
library(Hmisc)

## Loading required package: lattice
## Loading required package: survival
## Loading required package: Formula
##
## Attaching package: 'Hmisc'
##
## The following objects are masked from 'package:dplyr':
##
##     src, summarize
##
## The following objects are masked from 'package:base':
##
##     format.pval, units
```

Table 1: A table of the first few rows of the nycflights13 data.

| year.x | month | day | dep_time | sched_dep_time | dep_delay | arr_time | sched_arr_time | arr_delay | carrier | f |
|--------|-------|-----|----------|----------------|-----------|----------|----------------|-----------|---------|---|
| 2013 | 11 | 7 | 600 | 600 | 0 | 826 | 825 | 1 | WN | |
| 2013 | 10 | 30 | 1252 | 1250 | 2 | 1356 | 1400 | -4 | AA | |
| 2013 | 12 | 18 | 1723 | 1715 | 8 | 2008 | 2020 | -12 | DL | |
| 2013 | 11 | 20 | 2029 | 2030 | -1 | 2141 | 2205 | -24 | WN | |
| 2013 | 10 | 21 | 1620 | 1625 | -5 | 1818 | 1831 | -13 | DL | |
| 2013 | 11 | 7 | 852 | 900 | -8 | 1139 | 1157 | -18 | B6 | |

```

set.seed(42)
original_data <- read_csv("fltrain.csv.gz")

## Parsed with column specification:
## cols(
##   .default = col_double(),
##   carrier = col_character(),
##   tailnum = col_character(),
##   origin = col_character(),
##   dest = col_character(),
##   time_hour = col_datetime(format = ""),
##   name = col_character(),
##   dst = col_character(),
##   tzone = col_character(),
##   type = col_character(),
##   manufacturer = col_character(),
##   model = col_character(),
##   engine = col_character()
## )

## See spec(...) for full column specifications.
DF <- original_data

turning all columns with datatype characters to factors.
DF[sapply(DF, is.character)] <- lapply(DF[sapply(DF, is.character)],
                                       as.factor)
DF$flight <- as.factor(DF$flight)

library(lubridate)

##
## Attaching package: 'lubridate'
## The following object is masked from 'package:base':
##
##   date
DF$sched_arr_time_posix <- as.POSIXct(str_pad(as.character(DF$sched_arr_time), 4, pad="0"),format="%H%M")
DF$sched_arr_time_hour <- hour(DF$sched_arr_time_posix)
DF$sched_arr_time_minute <- minute(DF$sched_arr_time_posix)

#num minute is number of minutes since start of day for scheduled arrival time

```

```
DF$sched_arr_time_num_minute <- 60*DF$sched_arr_time_hour + DF$sched_arr_time_minute
```

```
DF$sched_dep_time_posix <- as.POSIXct(str_pad(as.character(DF$sched_dep_time),4 , pad="0"),format="%H%M",
```

```
DF$sched_dep_time_hour <- hour(DF$sched_dep_time_posix)
```

```
DF$sched_dep_time_minute <- minute(DF$sched_dep_time_posix)
```

```
#num minute is number of minutes since start of day for scheduled depival time
```

```
DF$sched_dep_time_num_minute <- 60*DF$sched_dep_time_hour + DF$sched_dep_time_minute
```

```
select(original_data, time_hour, sched_dep_time, sched_arr_time, tz, tzone)
```

```
## # A tibble: 200,000 x 5
```

```
##   time_hour      sched_dep_time sched_arr_time    tz tzone
##   <dtm>          <dbl>          <dbl> <dbl> <chr>
## 1 2013-11-07 11:00:00          600          825    -5 America/New_York
## 2 2013-10-30 16:00:00         1250         1400    -5 America/New_York
## 3 2013-12-18 22:00:00         1715         2020    -5 America/New_York
## 4 2013-11-21 01:00:00         2030         2205    -6 America/Chicago
## 5 2013-10-21 20:00:00         1625         1831    -5 America/New_York
## 6 2013-11-07 14:00:00          900         1157    -5 America/New_York
## 7 2013-09-29 19:00:00         1529         1649    -6 America/Chicago
## 8 2013-12-21 20:00:00         1530         1710    -6 America/Chicago
## 9 2013-11-07 21:00:00         1650         1906    -5 America/New_York
## 10 2013-03-31 21:00:00         1700         1821    -5 America/New_York
```

```
## # ... with 199,990 more rows
```

```
select(DF, sched_arr_time, sched_arr_time_hour)
```

```
## # A tibble: 200,000 x 2
```

```
##   sched_arr_time sched_arr_time_hour
##   <dbl>          <int>
## 1         825            8
## 2        1400           14
## 3        2020           20
## 4        2205           22
## 5        1831           18
## 6        1157           11
## 7        1649           16
## 8        1710           17
## 9        1906           19
## 10       1821           18
```

```
## # ... with 199,990 more rows
```

```
DF$sched_air_time <- DF$sched_arr_time_posix - DF$sched_dep_time_posix
```

```
drops <- c('sched_arr_time_posix', 'sched_arr_time_hour', 'sched_dep_time_posix', 'sched_dep_time_hour')
```

```
DF <- DF[ , !(names(DF) %in% drops)]
```

```
drops <- c("dep_time", "arr_time", "air_time", "arr_delay", "year.x", 'tailnum')
```

```
DF <- DF[ , !(names(DF) %in% drops)]
```

```
## Remove columns with more than 50% NA
```

```
DF <- DF[, -which(colMeans(is.na(DF)) > 0.5)]
```

```
DF$sched_air_time <- as.numeric(DF$sched_air_time)
```

```
library(imputeMissings)
```

```
##
```



```

## Attaching package: 'imputeMissings'
## The following object is masked from 'package:Hmisc':
##
##      impute
## The following object is masked from 'package:dplyr':
##
##      compute
impute_model <- imputeMissings::compute(DF, method="median/mode")
impute_model


## $month
## [1] 7
##
## $day
## [1] 16
##
## $dep_delay
## [1] -2
##
## $carrier
## [1] "UA"
##
## $flight
## [1] "15"
##
## $origin
## [1] "EWR"
##
## $dest
## [1] "ATL"
##
## $distance
## [1] 872
##
## $temp
## [1] 57.2
##
## $dewp
## [1] 42.8
##
## $humid
## [1] 57.69
##
## $wind_dir
## [1] 220
##
## $wind_speed
## [1] 10.35702
##
## $precip
## [1] 0
##
## $pressure

```

```

## [1] 1017.5
##
## $visib
## [1] 10
##
## $name
## [1] "Hartsfield Jackson Atlanta Intl"
##
## $lat
## [1] 36.09775
##
## $lon
## [1] -83.35339
##
## $alt
## [1] 433
##
## $tz
## [1] -5
##
## $dst
## [1] "A"
##
## $tzone
## [1] "America/New_York"
##
## $year.y
## [1] 2002
##
## $type
## [1] "Fixed wing multi engine"
##
## $manufacturer
## [1] "BOEING"
##
## $model
## [1] "A320-232"
##
## $engines
## [1] 2
##
## $seats
## [1] 149
##
## $engine
## [1] "Turbo-fan"
##
## $sched_arr_time_minute
## [1] 30
##
## $sched_arr_time_num_minute
## [1] 957
##
## $sched_dep_time_minute

```



```
## [1] 29
##
## $sched_dep_time_num_minute
## [1] 839
##
## $sched_air_time
## [1] 139

DF <- impute(DF, object=impute_model, flag=TRUE)
DF <- DF[!duplicated(as.list(DF))] #remove all redundant flag columns that are identical to each other

numeric_only_df <- dplyr::select_if(DF, is.numeric)
library(corrplot)

## corrplot 0.84 loaded
```

6.3 try features scaling

```
dep_delay_vec <- DF$dep_delay
DF$dep_delay <- NULL
head(DF)

##   month day carrier flight origin dest distance  temp  dewp humid wind_dir
## 1    11   7      WN   1716    LGA  ATL      762 62.96 55.94 77.83     210
```

```

## 2    10 30      AA    178    JFK BOS      187 59.00 46.94 64.22      240
## 3    12 18      DL    1585   LGA MCO      950 33.98 17.06 49.51      270
## 4    11 20      WN    3494   EWR MDW      711 37.04 17.96 45.58       20
## 5    10 21      DL    2231   LGA DTW      502 62.96 41.00 44.47      160
## 6    11 7       B6     27    EWR MCO      937 64.40 55.40 77.29      240
##   wind_speed precip pressure visib      name      lat
## 1    13.80936      0    1011.0    10   Hartsfield Jackson Atlanta Intl 33.63672
## 2     9.20624      0    1024.9    10 General Edward Lawrence Logan Intl 42.36435
## 3    17.26170      0    1019.8    10      Orlando Intl 28.42939
## 4     5.75390      0    1035.6    10      Chicago Midway Intl 41.78597
## 5    13.80936      0    1016.9    10      Detroit Metro Wayne Co 42.21244
## 6    16.11092      0    1017.5    10      Orlando Intl 28.42939
##      lon alt tz dst      tzone year.y      type
## 1 -84.42807 1026 -5   A America/New_York 2001 Fixed wing multi engine
## 2 -71.00518  19 -5   A America/New_York 2002 Fixed wing multi engine
## 3 -81.30899  96 -5   A America/New_York 2002 Fixed wing multi engine
## 4 -87.75242  620 -6   A America/Chicago 2006 Fixed wing multi engine
## 5 -83.35339  645 -5   A America/New_York 1992 Fixed wing multi engine
## 6 -81.30899  96 -5   A America/New_York 2006 Fixed wing multi engine
##      manufacturer      model engines seats      engine sched_arr_time_minute
## 1      BOEING 737-7H4      2    140 Turbo-fan      25
## 2      BOEING A320-232      2    149 Turbo-fan      0
## 3      AIRBUS A319-114      2    145 Turbo-fan      20
## 4      BOEING 737-7H4      2    140 Turbo-fan      5
## 5 AIRBUS INDUSTRIE A320-211      2    182 Turbo-jet     31
## 6      AIRBUS A320-232      2    200 Turbo-fan      57
##      sched_arr_time_num_minute sched_dep_time_minute sched_dep_time_num_minute
## 1      505      0      360
## 2      840     50     770
## 3     1220     15    1035
## 4     1325     30    1230
## 5     1111     25     985
## 6      717      0     540
##      sched_air_time dep_delay_flag temp_flag wind_dir_flag wind_speed_flag
## 1      145      0      0      0      0
## 2       70      0      0      0      0
## 3      185      0      0      0      0
## 4       95      0      0      0      0
## 5      126      0      0      0      0
## 6      177      0      0      0      0
##      precip_flag pressure_flag name_flag year.y_flag type_flag
## 1      0      0      0      0      0
## 2      0      0      0      1      1
## 3      0      0      0      0      0
## 4      0      0      0      0      0
## 5      0      0      0      0      0
## 6      0      1      0      0      0

```

```

library(dplyr)
DF <- DF %>% mutate_if(is.numeric, scale)
head(DF)

```

```

##      month      day carrier flight origin dest      distance      temp
## 1 1.30322 -0.9929373      WN   1716    LGA  ATL -0.3777852  0.3339858
## 2 1.01019  1.6325235      AA    178    JFK  BOS -1.1644742  0.1127815

```

```

## 3 1.59625 0.2627179 DL 1585 LGA MCO -0.1205721 -1.2848272
## 4 1.30322 0.4910188 WN 3494 EWR MDW -0.4475611 -1.1138966
## 5 1.01019 0.6051693 DL 2231 LGA DTW -0.7335054 0.3339858
## 6 1.30322 -0.9929373 B6 27 EWR MCO -0.1383581 0.4144237
##      dewp      humid      wind_dir wind_speed      precip      pressure      visib
## 1 0.7418623 0.9315583 0.07717317 0.4871566 -0.1492223 -0.96830167 0.3664282
## 2 0.2753242 0.2375566 0.36735789 -0.3415806 -0.1492223 1.01596006 0.3664282
## 3 -1.2735821 -0.5125364 0.65754261 1.1087096 -0.1492223 0.28792159 0.3664282
## 4 -1.2269283 -0.7129351 -1.76066338 -0.9631336 -0.1492223 2.54341334 0.3664282
## 5 -0.0325909 -0.7695363 -0.40646802 0.4871566 -0.1492223 -0.12606108 0.3664282
## 6 0.7138700 0.9040226 0.36735789 0.9015253 -0.1492223 -0.04040949 0.3664282
##      name      lat      lon      alt
## 1 Hartsfield Jackson Atlanta Intl -0.4207546 0.3298674 0.48364704
## 2 General Edward Lawrence Logan Intl 1.1190951 1.2378347 -0.60619417
## 3 Orlando Intl -1.3395034 0.5408516 -0.52285973
## 4 Chicago Midway Intl 1.0170501 0.1049977 0.04424731
## 5 Detroit Metro Wayne Co 1.0922942 0.4025621 0.07130394
## 6 Orlando Intl -1.3395034 0.5408516 -0.52285973
##      tz dst      tzone      year.y      type
## 1 0.6826595 A America/New_York -0.08500492 Fixed wing multi engine
## 2 0.6826595 A America/New_York 0.08617407 Fixed wing multi engine
## 3 0.6826595 A America/New_York 0.08617407 Fixed wing multi engine
## 4 -0.2514221 A America/Chicago 0.77089000 Fixed wing multi engine
## 5 0.6826595 A America/New_York -1.62561576 Fixed wing multi engine
## 6 0.6826595 A America/New_York 0.77089000 Fixed wing multi engine
##      manufacturer      model      engines      seats      engine
## 1 BOEING 737-7H4 0.05879311 0.02232546 Turbo-fan
## 2 BOEING A320-232 0.05879311 0.15869100 Turbo-fan
## 3 AIRBUS A319-114 0.05879311 0.09808410 Turbo-fan
## 4 BOEING 737-7H4 0.05879311 0.02232546 Turbo-fan
## 5 AIRBUS INDUSTRIE A320-211 0.05879311 0.65869797 Turbo-jet
## 6 AIRBUS A320-232 0.05879311 0.93142905 Turbo-fan
##      sched_arr_time_minute sched_arr_time_num_minute sched_dep_time_minute
## 1 -0.2348938 -1.4325947 -1.36042229
## 2 -1.6716145 -0.3129587 1.23408583
## 3 -0.5222379 0.9570761 -0.58206985
## 4 -1.3842703 1.3080068 0.19628258
## 5 0.1099192 0.5927766 -0.06316823
## 6 1.6041087 -0.7240489 -1.36042229
##      sched_dep_time_num_minute sched_air_time dep_delay_flag temp_flag
## 1 -1.6236293 0.14883297 0 0
## 2 -0.1673447 -0.24317221 0 0
## 3 0.7739125 0.35790240 0 0
## 4 1.4665357 -0.11250381 0 0
## 5 0.5963168 0.04952499 0 0
## 6 -0.9842849 0.31608852 0 0
##      wind_dir_flag wind_speed_flag precip_flag pressure_flag name_flag year.y_flag
## 1 0 0 0 0 0 0
## 2 0 0 0 0 0 1
## 3 0 0 0 0 0 0
## 4 0 0 0 0 0 0
## 5 0 0 0 0 0 0
## 6 0 0 0 1 0 0
##      type_flag

```

Table 2: proportion of variance explained by each principal component

| | x |
|-----|-----------|
| PC1 | 0.1352687 |
| PC2 | 0.1061946 |
| PC3 | 0.0862791 |
| PC4 | 0.0688860 |
| PC5 | 0.0605903 |
| PC6 | 0.0582950 |

```
## 1      0
## 2      1
## 3      0
## 4      0
## 5      0
## 6      0
```

```
DF$dep_delay <- dep_delay_vec
```

6.4 Exploratory Data Analysis

```
numeric_DF <- dplyr::select_if(DF, is.numeric) %>% scale()
```

```
prcomp_res <- prcomp(numeric_DF)
sdev <- prcomp_res$sdev
sdev
```

```
## [1] 1.801791e+00 1.596455e+00 1.438992e+00 1.285793e+00 1.205889e+00
## [6] 1.182827e+00 1.096174e+00 1.028253e+00 1.017278e+00 9.998650e-01
## [11] 9.870519e-01 9.598095e-01 9.527413e-01 9.302188e-01 8.722020e-01
## [16] 8.678655e-01 8.004108e-01 7.322880e-01 7.049102e-01 6.392628e-01
## [21] 2.137689e-01 1.418198e-01 5.881124e-02 1.734179e-14
```

6.4.1 all four components at same time

proportion of variance explained by each component

```
pve <- colSums(prcomp_res$x^2)/sum(numeric_DF^2)
```

```
rotation <- as.data.frame(prcomp_res$rotation)
rotation[order(-abs(rotation$PC1)),]
```

```
##           PC1           PC2           PC3           PC4
## lon      0.5342621313  0.008227987 -0.011812828 -0.0071811692
## distance -0.5303215501 -0.013672653  0.042091019  0.0006060091
## tz       0.5152320734  0.003918041 -0.003819370 -0.0093898478
## seats    -0.3267974610 -0.021806698  0.092397836  0.0075811906
## alt      -0.1291184773  0.001274092 -0.022430973  0.0179064904
## sched_air_time -0.1269319210 -0.059262301 -0.046756285 -0.0535120057
## year.y     0.1124548236  0.025774490 -0.038449639  0.0056654703
## sched_arr_time_num_minute -0.0808327463 -0.139515851 -0.548184284 -0.2068187847
```

| | | | | |
|------------------------------|---------------|--------------|--------------|---------------|
| ## lat | 0.0722096939 | 0.032190638 | -0.080234329 | 0.0246355126 |
| ## dewp | -0.0345819273 | 0.539843580 | -0.222885964 | 0.1989663082 |
| ## temp | -0.0291750632 | 0.384082292 | -0.336092763 | 0.3735114971 |
| ## month | -0.0237188202 | 0.182298459 | -0.084197143 | 0.3005556332 |
| ## sched_arr_time_minute | -0.0215092361 | -0.015025384 | -0.089818353 | -0.0429715263 |
| ## humid | -0.0203751852 | 0.453383850 | 0.144953536 | -0.3156006821 |
| ## engines | -0.0127603053 | 0.004264643 | 0.001902439 | 0.0110691777 |
| ## wind_speed | -0.0082656281 | -0.253349116 | -0.146083514 | -0.2029687361 |
| ## dep_delay | 0.0077853490 | 0.083641254 | -0.226497841 | -0.2377245680 |
| ## visib | 0.0071744619 | -0.280198624 | -0.146006623 | 0.5175332824 |
| ## pressure | 0.0051374481 | -0.127426965 | 0.233410657 | 0.1117576955 |
| ## precip | -0.0032312124 | 0.177264013 | 0.017625071 | -0.4019773885 |
| ## wind_dir | 0.0031602295 | -0.294478266 | -0.112869407 | 0.0491076591 |
| ## sched_dep_time_minute | -0.0026413644 | 0.003439200 | -0.121374246 | -0.0057857207 |
| ## day | 0.0004234420 | -0.029084254 | -0.002485255 | 0.0805143860 |
| ## sched_dep_time_num_minute | 0.0003538067 | -0.107997556 | -0.550807880 | -0.1834313382 |
| ## | PC5 | PC6 | PC7 | PC8 |
| ## lon | -0.141255373 | 0.021194537 | 0.014171904 | -0.023884973 |
| ## distance | -0.042941576 | -0.005942745 | -0.122735618 | 0.043705253 |
| ## tz | -0.207494095 | 0.030490047 | 0.001420389 | -0.017609684 |
| ## seats | -0.299154868 | -0.009581806 | -0.206563161 | -0.144397469 |
| ## alt | 0.466078117 | -0.062307363 | 0.241777658 | -0.082760758 |
| ## sched_air_time | -0.299112698 | 0.151635298 | 0.696202208 | -0.179246326 |
| ## year.y | 0.174145263 | 0.012033458 | -0.193371095 | -0.427597840 |
| ## sched_arr_time_num_minute | -0.104506952 | 0.310826773 | 0.201433129 | -0.047918183 |
| ## lat | 0.662475235 | -0.078116941 | 0.239929946 | -0.075262096 |
| ## dewp | -0.078261453 | -0.121386482 | 0.033430978 | 0.007641409 |
| ## temp | -0.092120074 | -0.132315189 | 0.049043959 | 0.010668654 |
| ## month | -0.018610656 | 0.151704760 | -0.001495157 | -0.086444244 |
| ## sched_arr_time_minute | -0.039500802 | 0.101328352 | -0.031997725 | -0.024890953 |
| ## humid | 0.003708481 | -0.016414325 | -0.017906587 | -0.005803837 |
| ## engines | -0.044314253 | 0.009660787 | -0.123433154 | -0.848006590 |
| ## wind_speed | -0.061901823 | -0.437801958 | 0.009361381 | -0.047218482 |
| ## dep_delay | 0.045376993 | -0.018945452 | -0.127408457 | 0.067633220 |
| ## visib | -0.007702427 | -0.018742792 | -0.026418970 | 0.024830862 |
| ## pressure | 0.116032120 | 0.563438391 | -0.052924000 | -0.001194143 |
| ## precip | -0.003734675 | -0.050856276 | 0.034962297 | -0.050209716 |
| ## wind_dir | -0.062792217 | -0.485144428 | 0.012771217 | -0.047966323 |
| ## sched_dep_time_minute | 0.092015290 | 0.054589781 | -0.390804984 | 0.021989590 |
| ## day | -0.001689151 | 0.020937857 | -0.047272003 | 0.037610315 |
| ## sched_dep_time_num_minute | 0.092202353 | 0.227284354 | -0.259042906 | 0.070884810 |
| ## | PC9 | PC10 | PC11 | PC12 |
| ## lon | -0.022482607 | 0.032770492 | -0.09242396 | 0.035571265 |
| ## distance | 0.020078457 | -0.041564282 | 0.11149965 | -0.041739208 |
| ## tz | -0.044272752 | 0.043751701 | -0.10878566 | 0.038322354 |
| ## seats | 0.015630688 | 0.063366532 | -0.24788381 | 0.086569086 |
| ## alt | -0.138469996 | 0.146822298 | -0.29987639 | 0.092299979 |
| ## sched_air_time | -0.149133954 | 0.035683038 | 0.19693563 | -0.143550228 |
| ## year.y | -0.034976443 | -0.155852739 | 0.66240241 | -0.262765846 |
| ## sched_arr_time_num_minute | 0.053926504 | -0.020110402 | 0.04077386 | -0.025711336 |
| ## lat | 0.017805635 | 0.016011358 | -0.03785986 | 0.012953704 |
| ## dewp | 0.001688937 | 0.004365652 | -0.01228182 | -0.051931274 |
| ## temp | 0.006200971 | -0.013446913 | -0.02987388 | -0.002801481 |
| ## month | -0.120787905 | -0.156460256 | 0.19154435 | 0.625378579 |

| | | | | |
|------------------------------|---------------|--------------|---------------|--------------|
| ## sched_arr_time_minute | -0.732058004 | 0.278259833 | -0.14647214 | -0.044943774 |
| ## humid | -0.010455120 | 0.037697965 | 0.03421562 | -0.091955092 |
| ## engines | 0.132086910 | 0.104896992 | -0.32454087 | 0.021879150 |
| ## wind_speed | -0.071625458 | -0.057666857 | 0.10354898 | 0.185290272 |
| ## dep_delay | 0.200325467 | 0.111865346 | -0.09362580 | -0.081177009 |
| ## visib | 0.032108130 | -0.048673404 | -0.05950128 | -0.052108675 |
| ## pressure | 0.010837763 | -0.093446997 | 0.03191945 | 0.236585784 |
| ## precip | -0.076608798 | -0.105843585 | 0.08898224 | 0.535791873 |
| ## wind_dir | -0.088916746 | -0.037067781 | 0.09755445 | 0.208042591 |
| ## sched_dep_time_minute | -0.502138130 | 0.043987210 | 0.09868480 | -0.123135660 |
| ## day | 0.213140360 | 0.886240079 | 0.33176453 | 0.184965679 |
| ## sched_dep_time_num_minute | 0.158656717 | -0.045621287 | -0.09049849 | 0.070227231 |
| ## | PC13 | PC14 | PC15 | PC16 |
| ## lon | -0.039446447 | -0.011082224 | -0.095304541 | 0.077597524 |
| ## distance | 0.039795147 | 0.002647274 | 0.039897734 | -0.004261933 |
| ## tz | -0.063720131 | -0.035598494 | -0.136768637 | 0.126354681 |
| ## seats | -0.027325189 | 0.020728319 | -0.252623268 | 0.247686911 |
| ## alt | -0.170421919 | -0.125274181 | -0.455475123 | 0.483583152 |
| ## sched_air_time | -0.164479367 | -0.208946245 | 0.041593996 | -0.030363222 |
| ## year.y | 0.197515002 | 0.011913488 | -0.161597312 | 0.354827113 |
| ## sched_arr_time_num_minute | -0.077917477 | 0.109167190 | -0.088599384 | -0.027228430 |
| ## lat | 0.035792204 | 0.086517293 | 0.199298954 | -0.284868167 |
| ## dewp | -0.017049314 | 0.049682409 | 0.061863302 | 0.060756471 |
| ## temp | -0.047848117 | 0.049638704 | 0.196084120 | 0.173608736 |
| ## month | 0.118816359 | -0.207210166 | -0.371210630 | -0.334318394 |
| ## sched_arr_time_minute | 0.556752791 | 0.125447499 | 0.111052124 | -0.001069467 |
| ## humid | 0.040272692 | 0.026173139 | -0.229179766 | -0.191945954 |
| ## engines | -0.057545879 | 0.002619518 | 0.178322651 | -0.213145500 |
| ## wind_speed | 0.003326375 | -0.006229440 | -0.077318196 | -0.076182657 |
| ## dep_delay | 0.290573141 | -0.828553616 | 0.116709581 | 0.042342062 |
| ## visib | 0.056744264 | -0.097207648 | 0.203736898 | 0.193394671 |
| ## pressure | 0.027199545 | -0.128666719 | 0.184057892 | 0.108434211 |
| ## precip | -0.102936363 | 0.077513423 | 0.479127840 | 0.404120615 |
| ## wind_dir | 0.042744908 | -0.080419991 | -0.006239399 | -0.069929842 |
| ## sched_dep_time_minute | -0.670765113 | -0.241104191 | 0.085424899 | -0.145499061 |
| ## day | -0.079973980 | 0.072516223 | 0.018126851 | 0.001947757 |
| ## sched_dep_time_num_minute | 0.028967986 | 0.258010080 | -0.122424720 | -0.008303141 |
| ## | PC17 | PC18 | PC19 | PC20 |
| ## lon | -0.0136295128 | 0.017750740 | 0.0968378420 | -0.013684012 |
| ## distance | 0.0205984456 | -0.007328044 | -0.1825459698 | 0.002448852 |
| ## tz | -0.0051770963 | 0.018809032 | 0.0334504287 | -0.017652697 |
| ## seats | -0.0560405964 | 0.016440667 | 0.7255605863 | 0.013581073 |
| ## alt | 0.0190843976 | 0.036647205 | -0.2601368362 | -0.018726070 |
| ## sched_air_time | -0.0034980154 | -0.004840517 | 0.0992775830 | -0.034805534 |
| ## year.y | 0.0096125145 | -0.009443594 | 0.0894488451 | 0.039095520 |
| ## sched_arr_time_num_minute | 0.0562858424 | 0.035350187 | 0.0356674787 | -0.080371931 |
| ## lat | -0.0520614544 | -0.012599070 | 0.5277297074 | -0.001035736 |
| ## dewp | -0.0942744599 | 0.231249850 | -0.0056935536 | -0.133117366 |
| ## temp | -0.1050475638 | 0.211514725 | -0.0068747953 | 0.258118643 |
| ## month | 0.0289369124 | -0.245815406 | -0.0116573597 | 0.055270714 |
| ## sched_arr_time_minute | -0.0117347846 | -0.001681325 | -0.0254211838 | 0.016161871 |
| ## humid | 0.0025024154 | 0.129312624 | -0.0006039130 | -0.656492439 |
| ## engines | 0.0086998591 | -0.003776375 | -0.2195595169 | -0.011263309 |
| ## wind_speed | -0.7670490408 | 0.118367572 | -0.0704672685 | -0.051948884 |

| | | | | |
|------------------------------|---------------|---------------|---------------|--------------|
| ## dep_delay | -0.0005171053 | -0.022901277 | 0.0568514930 | 0.061563276 |
| ## visib | -0.0735870801 | -0.324502693 | 0.0148202229 | -0.645786881 |
| ## pressure | -0.2197533179 | 0.626620074 | 0.0007539589 | -0.115345022 |
| ## precip | 0.1120149351 | -0.238971520 | 0.0196643203 | -0.093284192 |
| ## wind_dir | 0.5566040091 | 0.499155917 | 0.0416506799 | -0.142598451 |
| ## sched_dep_time_minute | 0.0036919766 | -0.018413525 | 0.0629338235 | -0.020942773 |
| ## day | -0.0242345214 | -0.001266925 | 0.0008713042 | -0.029416910 |
| ## sched_dep_time_num_minute | 0.0621948131 | 0.040857801 | -0.0295601569 | -0.061762457 |
| ## | PC21 | PC22 | PC23 | |
| ## lon | 0.051203785 | 0.8076315701 | 2.289417e-03 | |
| ## distance | -0.673700342 | 0.4367623709 | -5.904235e-04 | |
| ## tz | -0.692935877 | -0.3885478079 | -2.387247e-03 | |
| ## seats | 0.035356839 | -0.0141846176 | 3.536570e-05 | |
| ## alt | 0.004741932 | 0.0333354457 | 1.708179e-04 | |
| ## sched_air_time | -0.011013975 | 0.0180351813 | 6.012199e-04 | |
| ## year.y | 0.005825894 | -0.0149597542 | -4.767877e-04 | |
| ## sched_arr_time_num_minute | -0.006289146 | 0.0065552063 | -1.225185e-03 | |
| ## lat | -0.247687309 | 0.0630062661 | -3.594062e-04 | |
| ## dewp | 0.003122233 | 0.0006517912 | -7.090823e-01 | |
| ## temp | -0.005663154 | 0.0002352898 | 6.113164e-01 | |
| ## month | -0.001646308 | 0.0024681356 | 2.750428e-03 | |
| ## sched_arr_time_minute | 0.004516970 | 0.0014876945 | 2.250719e-05 | |
| ## humid | 0.009389784 | -0.0061721256 | 3.494391e-01 | |
| ## engines | -0.006154972 | 0.0044264994 | -3.645510e-05 | |
| ## wind_speed | 0.001061596 | -0.0025257566 | -5.755781e-04 | |
| ## dep_delay | -0.001497187 | 0.0024853904 | -1.169113e-03 | |
| ## visib | 0.003848279 | -0.0010054753 | 3.640126e-02 | |
| ## pressure | 0.002485028 | -0.0027882656 | 1.438635e-03 | |
| ## precip | -0.002132669 | 0.0014480599 | -5.757478e-03 | |
| ## wind_dir | 0.003375411 | -0.0014023794 | -3.540807e-04 | |
| ## sched_dep_time_minute | 0.017182271 | 0.0024772274 | -8.010616e-04 | |
| ## day | 0.001530860 | -0.0004536717 | -2.571908e-04 | |
| ## sched_dep_time_num_minute | 0.000800948 | -0.0052895688 | -1.710632e-03 | |
| ## | PC24 | | | |
| ## lon | 1.073158e-14 | | | |
| ## distance | 6.778678e-15 | | | |
| ## tz | -7.737005e-16 | | | |
| ## seats | 2.744135e-16 | | | |
| ## alt | -8.633487e-17 | | | |
| ## sched_air_time | -4.221606e-01 | | | |
| ## year.y | 4.918863e-16 | | | |
| ## sched_arr_time_num_minute | 6.602011e-01 | | | |
| ## lat | 2.167659e-15 | | | |
| ## dewp | -1.348152e-15 | | | |
| ## temp | -8.870884e-16 | | | |
| ## month | 2.656342e-16 | | | |
| ## sched_arr_time_minute | -1.997126e-15 | | | |
| ## humid | 2.401847e-15 | | | |
| ## engines | -1.067216e-16 | | | |
| ## wind_speed | -2.446632e-15 | | | |
| ## dep_delay | 5.015833e-17 | | | |
| ## visib | 5.676978e-16 | | | |
| ## pressure | 1.661713e-15 | | | |
| ## precip | -8.702010e-16 | | | |

Table 3: coefficients for each variable on each principal component

| | PC1 | PC2 | PC3 | PC4 | PC5 | PC6 | PC7 | PC8 |
|------------|------------|------------|------------|------------|------------|------------|------------|------------|
| dewp | -0.0345819 | 0.5398436 | -0.2228860 | 0.1989663 | -0.0782615 | -0.1213865 | 0.0334310 | 0.0076414 |
| humid | -0.0203752 | 0.4533838 | 0.1449535 | -0.3156007 | 0.0037085 | -0.0164143 | -0.0179066 | -0.0058038 |
| temp | -0.0291751 | 0.3840823 | -0.3360928 | 0.3735115 | -0.0921201 | -0.1323152 | 0.0490440 | 0.0106687 |
| wind_dir | 0.0031602 | -0.2944783 | -0.1128694 | 0.0491077 | -0.0627922 | -0.4851444 | 0.0127712 | -0.0479663 |
| visib | 0.0071745 | -0.2801986 | -0.1460066 | 0.5175333 | -0.0077024 | -0.0187428 | -0.0264190 | 0.0248309 |
| wind_speed | -0.0082656 | -0.2533491 | -0.1460835 | -0.2029687 | -0.0619018 | -0.4378020 | 0.0093614 | -0.0472185 |

```
## wind_dir                1.569069e-15
## sched_dep_time_minute   1.336574e-15
## day                    -3.738223e-16
## sched_dep_time_num_minute -6.212206e-01
pca_rotation <- head(rotation[order(-abs(rotation$PC2)),])
```

6.4.2 take out extreme departure delays

```
DF<-DF[DF$dep_delay < 30,]

set.seed(42)
DF$flight <- NULL
train_index <- sample(1:nrow(DF),size=2*nrow(DF)/3,replace=FALSE)
train_df <- DF[train_index,]
test_df <- DF[-train_index,]

# pre-allocate space
preallocate_df <- function(n){
  df <- data.frame(model_description = character(n), rmse = numeric(n), stringsAsFactors = FALSE)
  for(i in 1:n){
    df$model_description[i] <- i
    df$rmse[i] <- toString(i)
  }
  df
}
```

7 predicting 0

```
benchmark_df <- data.frame(model_description = character(), rmse = numeric(), stringsAsFactors = FALSE)
rmse = mean((test_df$dep_delay-0)^2) %>% sqrt()
model_description = "predicting 0"
benchmark_df <- rbind(benchmark_df, data.frame(model_description = model_description, rmse=rmse))
```

8 predicting the mean

```
rmse = mean((test_df$dep_delay - mean(train_df$dep_delay))^2) %>% sqrt()
rmse
```

```
## [1] 8.299767
```

```
benchmark_df
```

```
##   model_description   rmse
## 1      predicting 0 8.30571
```

9 predicting the median

```
rmse = mean((test_df$dep_delay - median(train_df$dep_delay))^2) %>% sqrt()
rmse
```

```
## [1] 8.469257
```

```
model_description <- 'predicting the median'
```

```
benchmark_df <- rbind(benchmark_df, data.frame(model_description = model_description, rmse=rmse))
```

10 linear regression with dep

```
model <- lm(dep_delay ~ ., data=train_df)
model_without_dep <- lm(dep_delay ~ .-dest, data=train_df)
anova(model, model_without_dep)
```

```
## Analysis of Variance Table
```

```
##
```

```
## Model 1: dep_delay ~ month + day + carrier + origin + dest + distance +
##   temp + dewp + humid + wind_dir + wind_speed + precip + pressure +
##   visib + name + lat + lon + alt + tz + dst + tzone + year.y +
##   type + manufacturer + model + engines + seats + engine +
##   sched_arr_time_minute + sched_arr_time_num_minute + sched_dep_time_minute +
##   sched_dep_time_num_minute + sched_air_time + dep_delay_flag +
##   temp_flag + wind_dir_flag + wind_speed_flag + precip_flag +
##   pressure_flag + name_flag + year.y_flag + type_flag
```

```
## Model 2: dep_delay ~ (month + day + carrier + origin + dest + distance +
##   temp + dewp + humid + wind_dir + wind_speed + precip + pressure +
##   visib + name + lat + lon + alt + tz + dst + tzone + year.y +
##   type + manufacturer + model + engines + seats + engine +
##   sched_arr_time_minute + sched_arr_time_num_minute + sched_dep_time_minute +
##   sched_dep_time_num_minute + sched_air_time + dep_delay_flag +
##   temp_flag + wind_dir_flag + wind_speed_flag + precip_flag +
##   pressure_flag + name_flag + year.y_flag + type_flag) - dest
```

```
##   Res.Df    RSS Df Sum of Sq    F    Pr(>F)
```

```
## 1 113488 7240975
```

```
## 2 113491 7242239 -3    -1263.7  6.602 0.0001863 ***
```

```
## ---
```

```
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
summary <- round(summary(model)$coefficients,6)
sorteddf <- summary[order(summary[,ncol(summary)]),]
head(sorteddf)

##           Estimate Std. Error    t value Pr(>|t|)
## carrierFL  4.615319   0.887510   5.200301     0
## carrierUA  2.440853   0.484369   5.039237     0
## originJFK  1.037568   0.117779   8.809419     0
## wind_speed 0.291278   0.027929  10.429081     0
## precip     0.206814   0.029257   7.068777     0
## pressure  -0.294549   0.027059 -10.885462     0
head(sorteddf)

##           Estimate Std. Error    t value Pr(>|t|)
## carrierFL  4.615319   0.887510   5.200301     0
## carrierUA  2.440853   0.484369   5.039237     0
## originJFK  1.037568   0.117779   8.809419     0
## wind_speed 0.291278   0.027929  10.429081     0
## precip     0.206814   0.029257   7.068777     0
## pressure  -0.294549   0.027059 -10.885462     0
lm_test_df <- test_df

in_test_but_not_train <- setdiff(unique(lm_test_df$model), unique(train_df$model))
lm_test_df <- lm_test_df[ !lm_test_df$model %in% in_test_but_not_train, ]

in_test_but_not_train <- setdiff(unique(lm_test_df$dest), unique(train_df$dest))
lm_test_df <- lm_test_df[ !lm_test_df$dest %in% in_test_but_not_train, ]

preds = predict(model, newdata=lm_test_df)

## Warning in predict.lm(model, newdata = lm_test_df): prediction from a rank-
## deficient fit may be misleading

rmse = sqrt(mean((lm_test_df$dep_delay - preds)^2))
rmse

## [1] 7.989994
model_description <- 'linear regression'
benchmark_df <- rbind(benchmark_df, data.frame(model_description = model_description, rmse=rmse))
```

11 gbm

```
library(gbm)

## Loaded gbm 2.1.5
train_gbm <- function(filename){
  set.seed(42)
  model <- gbm(dep_delay ~ ., data=train_df,
               n.trees=100, shrinkage=0.1) # default shrinkage = 0.1
  preds = predict(model, newdata=test_df, n.trees=100)
  rmse = sqrt(mean((test_df$dep_delay - preds)^2))
}
```


Table 4: gbm relative influence

| | var | rel.inf |
|---------------------------|---------------------------|-----------|
| sched_dep_time_num_minute | sched_dep_time_num_minute | 29.955850 |
| model | model | 26.206955 |
| dest | dest | 10.299706 |
| carrier | carrier | 8.907853 |
| month | month | 8.317570 |
| dewp | dewp | 4.563640 |

```
summary(model)
saveRDS(model, filename)
return(model)
}

destfile <- "models/model2.rds"
if (!file.exists(destfile)) {
  train_gbm(destfile)
}
model <- readRDS(destfile)

benchmark_df <- rbind(benchmark_df, data.frame(model_description = 'gbm', rmse=7.94458))
```



Here, you can see the relative influence for each variable for gbm.


For a gbm, the improvement in the splitting criterion (which is mean squared error for regression) for a given variable is calculated at each step. The relative influence for a given variable is the average of these improvements over all the trees where the aforementioned variable is used.

```
model <- gbm(dep_delay ~ ., data=train_df,
             n.trees=100, shrinkage=0.01) # default shrinkage = 0.1
```

```
## Distribution not specified, assuming gaussian ...
preds = predict(model, newdata=test_df, n.trees=1000)

## Warning in predict.gbm(model, newdata = test_df, n.trees = 1000): Number of
## trees not specified or exceeded number fit so far. Using 100.
filename <- "models/gbm_shrinkage_0point01_ntrees_100_v1.rds"
saveRDS(model, filename)
rmse = sqrt(mean((test_df$dep_delay - preds)^2))
rmse

## [1] 8.179232
summary(model)
```



```
##          var      rel.inf
## sched_dep_time_num_minute sched_dep_time_num_minute 55.4932927
## model          model 36.7164030
## carrier        carrier 6.4586978
## dest           dest 0.6914379
## sched_arr_time_num_minute sched_arr_time_num_minute 0.6401686
## month          month 0.0000000
## day            day 0.0000000
## origin         origin 0.0000000
## distance       distance 0.0000000
## temp           temp 0.0000000
## dewp           dewp 0.0000000
## humid          humid 0.0000000
## wind_dir       wind_dir 0.0000000
## wind_speed     wind_speed 0.0000000
## precip         precip 0.0000000
## pressure       pressure 0.0000000
## visib          visib 0.0000000
## name           name 0.0000000
```

```
## lat lat 0.0000000
## lon lon 0.0000000
## alt alt 0.0000000
## tz tz 0.0000000
## dst dst 0.0000000
## tzone tzone 0.0000000
## year.y year.y 0.0000000
## type type 0.0000000
## manufacturer manufacturer 0.0000000
## engines engines 0.0000000
## seats seats 0.0000000
## engine engine 0.0000000
## sched_arr_time_minute sched_arr_time_minute 0.0000000
## sched_dep_time_minute sched_dep_time_minute 0.0000000
## sched_air_time sched_air_time 0.0000000
## dep_delay_flag dep_delay_flag 0.0000000
## temp_flag temp_flag 0.0000000
## wind_dir_flag wind_dir_flag 0.0000000
## wind_speed_flag wind_speed_flag 0.0000000
## precip_flag precip_flag 0.0000000
## pressure_flag pressure_flag 0.0000000
## name_flag name_flag 0.0000000
## year.y_flag year.y_flag 0.0000000
## type_flag type_flag 0.0000000
```


```
library(gbm)
filename <- "models/gbm_shrinkage_0point01_ntrees_100_v1.rds"
if (!file.exists(filename)) {
  model <- gbm(
    dep_delay ~ .,
    data = train_df,
    n.trees = 100,
    shrinkage = 0.01
  ) # default shrinkage = 0.1
  saveRDS(model, filename)
} else {
  print("reading saved model")
  model <- readRDS(filename)
}
```

```
## [1] "reading saved model"
```

```
preds = predict(model, newdata = test_df, n.trees = 100)
rmse = sqrt(mean((test_df$dep_delay - preds) ^ 2))
rmse
```

```
## [1] 8.179232
```

```
summary(model)
```



| ## | | var | rel.inf |
|----|---------------------------|---------------------------|------------|
| ## | sched_dep_time_num_minute | sched_dep_time_num_minute | 55.4932927 |
| ## | model | model | 36.7164030 |
| ## | carrier | carrier | 6.4586978 |
| ## | dest | dest | 0.6914379 |
| ## | sched_arr_time_num_minute | sched_arr_time_num_minute | 0.6401686 |
| ## | month | month | 0.0000000 |
| ## | day | day | 0.0000000 |
| ## | origin | origin | 0.0000000 |
| ## | distance | distance | 0.0000000 |
| ## | temp | temp | 0.0000000 |
| ## | dewp | dewp | 0.0000000 |
| ## | humid | humid | 0.0000000 |
| ## | wind_dir | wind_dir | 0.0000000 |
| ## | wind_speed | wind_speed | 0.0000000 |
| ## | precip | precip | 0.0000000 |
| ## | pressure | pressure | 0.0000000 |
| ## | visib | visib | 0.0000000 |
| ## | name | name | 0.0000000 |
| ## | lat | lat | 0.0000000 |
| ## | lon | lon | 0.0000000 |
| ## | alt | alt | 0.0000000 |
| ## | tz | tz | 0.0000000 |
| ## | dst | dst | 0.0000000 |
| ## | tzone | tzone | 0.0000000 |
| ## | year.y | year.y | 0.0000000 |
| ## | type | type | 0.0000000 |
| ## | manufacturer | manufacturer | 0.0000000 |
| ## | engines | engines | 0.0000000 |
| ## | seats | seats | 0.0000000 |
| ## | engine | engine | 0.0000000 |
| ## | sched_arr_time_minute | sched_arr_time_minute | 0.0000000 |
| ## | sched_dep_time_minute | sched_dep_time_minute | 0.0000000 |


```
## sched_air_time          sched_air_time 0.0000000
## dep_delay_flag          dep_delay_flag 0.0000000
## temp_flag               temp_flag      0.0000000
## wind_dir_flag           wind_dir_flag   0.0000000
## wind_speed_flag         wind_speed_flag 0.0000000
## precip_flag             precip_flag     0.0000000
## pressure_flag           pressure_flag   0.0000000
## name_flag               name_flag       0.0000000
## year.y_flag             year.y_flag     0.0000000
## type_flag               type_flag       0.0000000
```

```
rmse = sqrt(mean((test_df$dep_delay - preds)^2))
rmse
```

```
## [1] 8.179232
```

```
set.seed(42)
x <- 2 ^ seq(5, 14, by = 1)
rmse_vec <- numeric(length(x))
count <- 1
for (val in x) {
  filename <-
    paste0("models/gbm_shrinkage_0point001_ntrees_", val)
  filename <- paste0(filename, "_v1.rds")
  if (!file.exists(filename)) {
    hboost <- gbm(
      dep_delay ~ .,
      data = train_df,
      n.trees = val,
      distribution = 'gaussian',
      shrinkage = 0.001
    )
    saveRDS(hboost, filename)
  } else{
    hboost <- readRDS(filename)
  }

  preds = predict(hboost, n.trees = val, newdata = test_df)
  mse = mean((test_df$dep_delay - preds) ^ 2)
  rmse <- sqrt(mse)
  rmse_vec[count] <- rmse
  print(val)
  print(rmse)
  count = count + 1
}
```

```
## [1] 32
## [1] 8.247971
## [1] 64
## [1] 8.211231
## [1] 128
## [1] 8.159652
## [1] 256
## [1] 8.101211
## [1] 512
## [1] 8.038247
```

```
## Warning in predict.gbm(hboost, n.trees = val, newdata = test_df): Number of
## trees not specified or exceeded number fit so far. Using 1000.

## [1] 1024
## [1] 7.980282

## Warning in predict.gbm(hboost, n.trees = val, newdata = test_df): Number of
## trees not specified or exceeded number fit so far. Using 1000.

## [1] 2048
## [1] 7.980282

## Warning in predict.gbm(hboost, n.trees = val, newdata = test_df): Number of
## trees not specified or exceeded number fit so far. Using 1000.

## [1] 4096
## [1] 7.980282


## Warning in predict.gbm(hboost, n.trees = val, newdata = test_df): Number of
## trees not specified or exceeded number fit so far. Using 1000.

## [1] 8192
## [1] 7.980282

## Warning in predict.gbm(hboost, n.trees = val, newdata = test_df): Number of
## trees not specified or exceeded number fit so far. Using 1000.

## [1] 16384
## [1] 7.980282
```

```
summary(hboost)
```



```
##          var      rel.inf
## sched_dep_time_num_minute sched_dep_time_num_minute 30.88424567
## model              model 25.67333688
## dest              dest 10.06233500
## carrier            carrier 8.92374891
## month              month 8.19668822
```

```

## dewp                dewp  4.37183782
## origin              origin 3.10232960
## sched_arr_time_num_minute sched_arr_time_num_minute 2.61899718
## precip              precip 2.32288426
## pressure            pressure 1.50793482
## humid              humid 1.10879519
## dep_delay_flag      dep_delay_flag 0.90034697
## temp              temp 0.12324130
## pressure_flag      pressure_flag 0.11994578
## day              day 0.05664389
## sched_dep_time_minute sched_dep_time_minute 0.02668850
## distance          distance 0.00000000
## wind_dir          wind_dir 0.00000000
## wind_speed        wind_speed 0.00000000
## visib            visib 0.00000000
## name            name 0.00000000
## lat            lat 0.00000000
## lon            lon 0.00000000
## alt            alt 0.00000000
## tz            tz 0.00000000
## dst            dst 0.00000000
## tzone          tzone 0.00000000
## year.y          year.y 0.00000000
## type            type 0.00000000
## manufacturer    manufacturer 0.00000000
## engines          engines 0.00000000
## seats           seats 0.00000000
## engine          engine 0.00000000
## sched_arr_time_minute sched_arr_time_minute 0.00000000
## sched_air_time    sched_air_time 0.00000000
## temp_flag        temp_flag 0.00000000
## wind_dir_flag    wind_dir_flag 0.00000000
## wind_speed_flag  wind_speed_flag 0.00000000
## precip_flag      precip_flag 0.00000000
## name_flag        name_flag 0.00000000
## year.y_flag      year.y_flag 0.00000000
## type_flag        type_flag 0.00000000

```

```


write.csv(summary,
            'performance/gbm_shrinkage_0point001_16384trees_summary_v1.csv')

plot(x, rmse_vec)

```

Table 5: benchmark comparing all the models that we tried

| model_description | rmse |
|-----------------------|----------|
| predicting 0 | 8.305710 |
| predicting the median | 8.469257 |
| linear regression | 7.989994 |
| gbm | 7.944580 |



```
num_trees_vs_rmse <- data.frame("num_trees" = x, "rmse" = rmse_vec)
write.csv(
  num_trees_vs_rmse,
  'performance/gbm_shrinkage_0point001_num_trees_vs_rmse_v1.csv'
)
```

References

“Gradient Boosting Machines · UC Business Analytics R Programming Guide.” 2019. http://uc-r.github.io/gbm_regression#h2o.