

knit

Vincent Chiu

4/1/2019

Contents

1	Introduction	2
2	Data	2
3	Methods:	2
3.1	Data Preprocessing	2
3.2	Modelling	2
3.3	Basic Models	2
3.4	Linear Regression	2
3.5	GBM	3
4	Results	3
5	Conclusion and Discussion	3
6	have train, CV and test set	4
7	1/3 train, 1/3 CV, 1/3 test	4
8	2/3% train, 1/3%CV, wait for prof test set	4
9	Code	4
9.1	Loading Libraries	4
9.2	Loading the data	5
10	turning all columns with datatype characters to factors.	6
11	Methods	8
11.1	Preprocessing	8
11.2	- Dropping columns that contain data from after the planes' departure which may leak information about the response variable dep_delay.	8
12	try features scaling	12
13	predicting 0	15
14	predicting the mean	15
15	predicting the median	15
16	linear regression with dep	16
17	gbm	17
	References	22
1.	Introduction (brief)	
2.	Data (brief)	

- 3. Methods
- 4. Results
- 5. Conclusions and Discussion

1 Introduction

The goal of this project is to predict the response variable, departure delays for a particular flight given the explanatory variables.

2 Data

The dataset consists of information about all the flights leaving from New York City in 2013. The dataset contains 43 variables in total. The dataset is an amalgamation of several datasets including datasets containing information on weather, the airports, the flights, and the models of airplanes. The training dataset provided to us contains 200,000 observations.

3 Methods:

3.1 Data Preprocessing

I performed data preprocessing. My data preprocessing steps include the following: 1. Dropping columns that contain data from after the planes' departure which may leak information about the response variable `dep_delay`. 2. Dropping columns with too many NAs. 3. Impute NAs for the remaining columns. 4. Scaling the data to work well with methods like lasso regression. 5. Only kept data which had a departure delay of less than 30 minutes late, which reduced the dataset from 200,000 rows to approximately 170,000.

3.2 Modelling

Initially, I used the most basic cross validation technique where I have a training dataset and a holdout test dataset. I split the original data into a ratio of 2/3 train and 1/3 of the data for test. I believe that this split gives enough data for the models to learn while 1/3 is enough data for me to get an accurate assessment of the error. k-folds cross validation was not initially used in order to save on compute time as I was only exploring the models. k-folds cross validation would increase training time for the models by a factor of k.

3.3 Basic Models

`dep_delay` is the number of minutes that the plane either departs early or late. Negative numbers are for early departures and positive numbers are for the number of minutes the plane is late. First, I used a basic model of simply predicting the `dep_delay` to always be 0. This was done to establish baseline performance. This model had an root mean squared error (RMSE) of 8.30571. TODO The model in which I predicted the mean for all the predictions had an RMSE of TODO.

3.4 Linear Regression

Then I tried linear regression with `dep_delay` as the response variables and all the other remaining variables as the explanatory variables. This model was better than predicting the mean with an RMSE of TODO. This suggests that there is some relationship between the `dep_delay` and the explanatory variables.

3.5 GBM

Aftwards, I tried a Generalized Boosted Regression Model (GBM). This model had the lowest RMSE on the test dataset after I tuned it to have a shrinkage of 0.01 and around 16,000 trees. Shrinkage is proportional to the learning rate. 16,000 trees is the number of trees used in the model. Each iteration uses 1 tree, so 16,000 trees also refers to the number of iterations. According to the vignette, the rmse can always be improved by decreasing shrinkage but this provides diminishing returns. A good strategy would be to pick a small shrinkage that balances performance and compute time. Then with this fixed shrinkage value, increase the number of trees until you get diminishing returns.

4 Results

In regression and gbm, I found different features to be important. For the best gbm model, dest which refers to which airport a given plane is going to was the most important feature. However, the one hot encoding versions of carrier were the most important features for regression. On the other hand, dest does appear as an important feature in linear regression as well but it is not the most important feature. I surmise that if we can somehow sum up all the contributions from each of the one hot encoded variables derived from dest then, it might appear as the most important feature for linear regression as well. We can try Anova in order to measure the statistical significance of dest. Performing anova on comparing linear regression model with and without dest, it was determined that due to the low p-value of 0.0001863 associated with having dest that keeping at least one of the one hot categorical variables derived from dest is beneficial for the linear regression model.

TODO: try interaction terms , try anova.

5 Conclusion and Discussion

Conclusion: In conclusion, out of the methods that we covered in class, I found gradient boosted models to provide the best performance based on having the lowest root mean squared error on the hold out test set.

Based on the relative influence scores provided by the gbm, some of the most important feature variables include dest, model, and sched_dep_time_num_minute.

The dest column contains the airport code for where a given flight is flying to. Based on my run of gbm with a shrinkage of 0.01 and 16834 trees, dest was the most important feature with 49.56 relative influence. (“Gradient Boosting Machines · UC Business Analytics R Programming Guide” 2019).

TODO; think about removing points that are outliers aka points with high cook’s distance consider removing outliers in train but not in test, then use k-folds cross validation on test.

TODO: remove points that are outliers ie dep_delay > 200 or 300 etc. or remove less than x number of points. then use k-folds cross validation on cross validation set where no points were removed. can repeat k-folds for different seeds. can just try this on my quickest model, ie linear regression. should be bowl shape vs rmse vs. number of points removed. theoretically

I also considered removing based on cook’s distance but this took too long to compute.

5 folds with 10 different random seeds

6 have train, CV and test set

7 1/3 train, 1/3 CV, 1/3 test

8 2/3% train, 1/3%CV, wait for prof test set

try lasso regression

9 Code

9.1 Loading Libraries

```
library(tidyverse)
```

```
## -- Attaching packages ----- tidyverse 1.2.1
```

```
## v ggplot2 3.2.1    v purrr  0.3.3
```

```
## v tibble  2.1.3    v dplyr  0.8.3
```

```
## v tidyr   1.0.0    v stringr 1.4.0
```

```
## v readr   1.3.1    v forcats 0.4.0
```

```
## -- Conflicts ----- tidyverse_conflicts()
```

```
## x dplyr::filter() masks stats::filter()
```

```
## x dplyr::lag()    masks stats::lag()
```


```
attach(mtcars)
```

```
## The following object is masked from package:ggplot2:
```

```
##
```

```
##      mpg
```

```
plot(wt, mpg)
```



9.2 Loading the data

```
library(nycflights13)
library(Hmisc)

## Loading required package: lattice
## Loading required package: survival
## Loading required package: Formula
##
## Attaching package: 'Hmisc'
## The following objects are masked from 'package:dplyr':
##
##   src, summarize
## The following objects are masked from 'package:base':
##
##   format.pval, units

set.seed(42)
original_data <- read_csv("fltrain.csv.gz")

## Parsed with column specification:
## cols(
##   .default = col_double(),
##   carrier = col_character(),
##   tailnum = col_character(),
##   origin = col_character(),
##   dest = col_character(),
##   time_hour = col_datetime(format = ""),
##   name = col_character(),
```

```
##   dst = col_character(),
##   tzone = col_character(),
##   type = col_character(),
##   manufacturer = col_character(),
##   model = col_character(),
##   engine = col_character()
## )

## See spec(...) for full column specifications.
DF <- original_data
```

10 turning all columns with datatype characters to factors.

```
DF[sapply(DF, is.character)] <- lapply(DF[sapply(DF, is.character)],
                                       as.factor)
DF$flight <- as.factor(DF$flight)
str(DF)

## Classes 'spec_tbl_df', 'tbl_df', 'tbl' and 'data.frame': 200000 obs. of  43 variables:
##  $ year.x      : num  2013 2013 2013 2013 2013 ...
##  $ month       : num  11 10 12 11 10 11 9 12 11 3 ...
##  $ day         : num  7 30 18 20 21 7 29 21 7 31 ...
##  $ dep_time    : num  600 1252 1723 2029 1620 ...
##  $ sched_dep_time: num  600 1250 1715 2030 1625 ...
##  $ dep_delay   : num  0 2 8 -1 -5 -8 -10 -4 0 -8 ...
##  $ arr_time    : num  826 1356 2008 2141 1818 ...
##  $ sched_arr_time: num  825 1400 2020 2205 1831 ...
##  $ arr_delay   : num  1 -4 -12 -24 -13 -18 -10 -16 4 -11 ...
##  $ carrier     : Factor w/ 16 levels "9E","AA","AS",...: 15 2 5 15 5 4 6 6 1 13 ...
##  $ flight      : Factor w/ 3672 levels "1","2","3","4",...: 1525 147 1400 2343 1860 24 3083 3351 20 ...
##  $ tailnum     : Factor w/ 3957 levels "D942DN","NOEGMQ",...: 1437 1226 836 565 756 2459 204 2890 6 ...
##  $ origin      : Factor w/ 3 levels "EWR","JFK","LGA": 3 2 3 1 3 1 1 3 2 3 ...
##  $ dest        : Factor w/ 104 levels "ABQ","ACK","ALB",...: 5 12 54 55 33 54 59 59 27 29 ...
##  $ air_time    : num  123 44 133 107 90 136 110 118 101 47 ...
##  $ distance    : num  762 187 950 711 502 937 725 738 589 214 ...
##  $ hour        : num  6 12 17 20 16 9 15 15 16 17 ...
##  $ minute      : num  0 50 15 30 25 0 29 30 50 0 ...
##  $ time_hour   : POSIXct, format: "2013-11-07 11:00:00" "2013-10-30 16:00:00" ...
##  $ temp        : num  63 59 34 37 63 ...
##  $ dewp        : num  55.9 46.9 17.1 18 41 ...
##  $ humid       : num  77.8 64.2 49.5 45.6 44.5 ...
##  $ wind_dir    : num  210 240 270 20 160 240 180 190 320 140 ...
##  $ wind_speed  : num  13.81 9.21 17.26 5.75 13.81 ...
##  $ wind_gust   : num  NA NA 21.9 NA NA ...
##  $ precip      : num  0 0 0 0 0 0 0 0 0 0 ...
##  $ pressure    : num  1011 1025 1020 1036 1017 ...
##  $ visib       : num  10 10 10 10 10 10 10 10 10 10 ...
##  $ name        : Factor w/ 100 levels "Akron Canton Regional Airport",...: 37 31 67 17 26 67 32 32 ...
##  $ lat         : num  33.6 42.4 28.4 41.8 42.2 ...
##  $ lon         : num  -84.4 -71 -81.3 -87.8 -83.4 ...
##  $ alt         : num  1026 19 96 620 645 ...
##  $ tz          : num  -5 -5 -5 -6 -5 -5 -6 -6 -5 -5 ...
```

```

## $ dst          : Factor w/ 2 levels "A","N": 1 1 1 1 1 1 1 1 1 1 ...
## $ tzone        : Factor w/ 7 levels "America/Anchorage",...: 5 5 5 2 5 5 2 2 5 5 ...
## $ year.y       : num 2001 NA 2002 2006 1992 ...
## $ type         : Factor w/ 3 levels "Fixed wing multi engine",...: 1 NA 1 1 1 1 1 1 1 1 ...
## $ manufacturer : Factor w/ 35 levels "AGUSTA SPA","AIRBUS",...: 10 NA 2 10 3 2 18 11 11 3 ...
## $ model        : Factor w/ 126 levels "150","172E","172M",...: 37 NA 80 37 84 88 106 98 99 79 ...
## $ engines      : num 2 NA 2 2 2 2 2 2 2 2 ...
## $ seats        : num 140 NA 145 140 182 200 55 80 95 179 ...
## $ speed        : num NA NA NA NA NA NA NA NA NA NA ...
## $ engine       : Factor w/ 6 levels "4 Cycle","Reciprocating",...: 3 NA 3 3 4 3 3 3 3 3 ...
## - attr(*, "spec")=
## .. cols(
## ..   year.x = col_double(),
## ..   month = col_double(),
## ..   day = col_double(),
## ..   dep_time = col_double(),
## ..   sched_dep_time = col_double(),
## ..   dep_delay = col_double(),
## ..   arr_time = col_double(),
## ..   sched_arr_time = col_double(),
## ..   arr_delay = col_double(),
## ..   carrier = col_character(),
## ..   flight = col_double(),
## ..   tailnum = col_character(),
## ..   origin = col_character(),
## ..   dest = col_character(),
## ..   air_time = col_double(),
## ..   distance = col_double(),
## ..   hour = col_double(),
## ..   minute = col_double(),
## ..   time_hour = col_datetime(format = ""),
## ..   temp = col_double(),
## ..   dewp = col_double(),
## ..   humid = col_double(),
## ..   wind_dir = col_double(),
## ..   wind_speed = col_double(),
## ..   wind_gust = col_double(),
## ..   precip = col_double(),
## ..   pressure = col_double(),
## ..   visib = col_double(),
## ..   name = col_character(),
## ..   lat = col_double(),
## ..   lon = col_double(),
## ..   alt = col_double(),
## ..   tz = col_double(),
## ..   dst = col_character(),
## ..   tzone = col_character(),
## ..   year.y = col_double(),
## ..   type = col_character(),
## ..   manufacturer = col_character(),
## ..   model = col_character(),
## ..   engines = col_double(),
## ..   seats = col_double(),
## ..   speed = col_double(),

```

```
## .. engine = col_character()
## .. )
```

11 Methods

11.1 Preprocessing

Data preprocessing steps include the following: - Dropping columns that contain data from after the planes' departure which may leak information about the response variable dep_delay. - Dropping columns with too many NAs. - Impute NAs for the remaining columns. - Scaling the data to work well with methods like lasso regression.

11.2 - Dropping columns that contain data from after the planes' departure which may leak information about the response variable dep_delay.

dropping the columns "dep_time", "arr_time", "air_time", "arr_delay", because that leaks the response variable. dropping column "year.x" because all the values are 2013 dropping tailnum because it produces too many dummy variable columns for one hot encoding.

```
library(lubridate)
```

```
##
## Attaching package: 'lubridate'
```

```
## The following object is masked from 'package:base':
```

```
##
```

```
## date
```

```
DF$sched_arr_time_posix <- as.POSIXct(str_pad(as.character(DF$sched_arr_time), 4, pad="0"),format="%H%M")
DF$sched_arr_time_hour <- hour(DF$sched_arr_time_posix)
DF$sched_arr_time_minute <- minute(DF$sched_arr_time_posix)
```

```
#num minute is number of minutes since start of day for scheduled arrival time
DF$sched_arr_time_num_minute <- 60*DF$sched_arr_time_hour + DF$sched_arr_time_minute
```

```
DF$sched_dep_time_posix <- as.POSIXct(str_pad(as.character(DF$sched_dep_time),4 , pad="0"),format="%H%M")
DF$sched_dep_time_hour <- hour(DF$sched_dep_time_posix)
DF$sched_dep_time_minute <- minute(DF$sched_dep_time_posix)
```

```
#num minute is number of minutes since start of day for scheduled depival time
DF$sched_dep_time_num_minute <- 60*DF$sched_dep_time_hour + DF$sched_dep_time_minute
```

```
select(original_data, time_hour, sched_dep_time, sched_arr_time, tz, tzone)
```

```
## # A tibble: 200,000 x 5
```

```
##   time_hour      sched_dep_time sched_arr_time    tz tzone
##   <dtm>          <dbl>          <dbl> <dbl> <chr>
## 1 2013-11-07 11:00:00          600           825    -5 America/New_York
## 2 2013-10-30 16:00:00         1250          1400    -5 America/New_York
## 3 2013-12-18 22:00:00         1715          2020    -5 America/New_York
## 4 2013-11-21 01:00:00         2030          2205    -6 America/Chicago
## 5 2013-10-21 20:00:00         1625          1831    -5 America/New_York
## 6 2013-11-07 14:00:00          900          1157    -5 America/New_York
## 7 2013-09-29 19:00:00         1529          1649    -6 America/Chicago
```



```
## 8 2013-12-21 20:00:00      1530      1710      -6 America/Chicago
## 9 2013-11-07 21:00:00      1650      1906      -5 America/New_York
## 10 2013-03-31 21:00:00      1700      1821      -5 America/New_York
## # ... with 199,990 more rows
```

```
select(DF, sched_arr_time, sched_arr_time_hour)
```

```
## # A tibble: 200,000 x 2
##   sched_arr_time sched_arr_time_hour
##   <dbl>          <int>
## 1         825           8
## 2        1400          14
## 3        2020          20
## 4        2205          22
## 5        1831          18
## 6        1157          11
## 7        1649          16
## 8        1710          17
## 9        1906          19
## 10       1821          18
## # ... with 199,990 more rows
```

```
DF$sched_air_time <- DF$sched_arr_time_posix - DF$sched_dep_time_posix
drops <- c('sched_arr_time_posix', 'sched_arr_time_hour', 'sched_dep_time_posix', 'sched_dep_time_hour')
DF <- DF[, !(names(DF) %in% drops)]
```

```
drops <- c("dep_time", "arr_time", "air_time", "arr_delay", "year.x", 'tailnum')
DF <- DF[, !(names(DF) %in% drops)]
```

```
DF
```

```
## # A tibble: 200,000 x 37
##   month   day dep_delay carrier flight origin dest distance temp dewp humid
##   <dbl> <dbl>   <dbl> <fct>   <fct> <fct> <fct>   <dbl> <dbl> <dbl> <dbl>
## 1    11     7         0 WN      1716 LGA    ATL     762  63.0  55.9  77.8
## 2    10    30         2 AA       178 JFK    BOS     187  59    46.9  64.2
## 3    12    18         8 DL      1585 LGA    MCO     950  34.0  17.1  49.5
## 4    11    20        -1 WN      3494 EWR    MDW     711  37.0  18.0  45.6
## 5    10    21        -5 DL      2231 LGA    DTW     502  63.0  41    44.5
## 6    11     7        -8 B6       27   EWR    MCO     937  64.4  55.4  77.3
## 7     9    29       -10 EV      4580 EWR    MKE     725  69.1  53.1  56.7
## 8    12    21        -4 EV      5207 LGA    MKE     738  57.9  46.0  64.5
## 9    11     7         0 9E      2910 JFK    CVG     589  53.6  48.2  81.9
## 10   3    31        -8 US      2183 LGA    DCA     214  51.1  36.0  56.0
## # ... with 199,990 more rows, and 26 more variables: wind_dir <dbl>,
## #   wind_speed <dbl>, wind_gust <dbl>, precip <dbl>, pressure <dbl>,
## #   visib <dbl>, name <fct>, lat <dbl>, lon <dbl>, alt <dbl>, tz <dbl>,
## #   dst <fct>, tzone <fct>, year.y <dbl>, type <fct>, manufacturer <fct>,
## #   model <fct>, engines <dbl>, seats <dbl>, speed <dbl>, engine <fct>,
## #   sched_arr_time_minute <int>, sched_arr_time_num_minute <dbl>,
## #   sched_dep_time_minute <int>, sched_dep_time_num_minute <dbl>,
## #   sched_air_time <drtn>
```

```
## Remove columns with more than 50% NA
DF <- DF[, -which(colMeans(is.na(DF)) > 0.5)]
```

```

DF$sched_air_time <- as.numeric(DF$sched_air_time)
library(imputeMissings)

##
## Attaching package: 'imputeMissings'
## The following object is masked from 'package:Hmisc':
##
##      impute
## The following object is masked from 'package:dplyr':
##
##      compute
impute_model <- imputeMissings::compute(DF, method="median/mode")
impute_model

## $month
## [1] 7
##
## $day
## [1] 16
##
## $dep_delay
## [1] -2
##
## $carrier
## [1] "UA"
##
## $flight
## [1] "15"
##
## $origin
## [1] "EWR"
##
## $dest
## [1] "ATL"
##
## $distance
## [1] 872
##
## $temp
## [1] 57.2
##
## $dewp
## [1] 42.8
##
## $humid
## [1] 57.69
##
## $wind_dir
## [1] 220
##
## $wind_speed
## [1] 10.35702

```

```

##
## $precip
## [1] 0
##
## $pressure
## [1] 1017.5
##
## $visib
## [1] 10
##
## $name
## [1] "Hartsfield Jackson Atlanta Intl"
##
## $lat
## [1] 36.09775
##
## $lon
## [1] -83.35339
##
## $alt
## [1] 433
##
## $tz
## [1] -5
##
## $dst
## [1] "A"
##
## $tzone
## [1] "America/New_York"
##
## $year.y
## [1] 2002
##
## $type
## [1] "Fixed wing multi engine"
##
## $manufacturer
## [1] "BOEING"
##
## $model
## [1] "A320-232"
##
## $engines
## [1] 2
##
## $seats
## [1] 149
##
## $engine
## [1] "Turbo-fan"
##
## $sched_arr_time_minute
## [1] 30

```

```
##
## $sched_arr_time_num_minute
## [1] 957
##
## $sched_dep_time_minute
## [1] 29
##
## $sched_dep_time_num_minute
## [1] 839
##
## $sched_air_time
## [1] 139


DF <- impute(DF, object=impute_model, flag=TRUE)
DF <- DF[!duplicated(as.list(DF))]
```

#remove all redundant flag columns that are identical to each other

```
numeric_only_df <- dplyr::select_if(DF, is.numeric)
library(corrplot)

## corrplot 0.84 loaded

corrplot(cor(numeric_only_df), type = 'lower')
```



12 try features scaling

```

dep_delay_vec <- DF$dep_delay
DF$dep_delay <- NULL
head(DF)

```

```

##   month day carrier flight origin dest distance  temp  dewp humid wind_dir
## 1    11   7      WN   1716   LGA  ATL       762 62.96 55.94 77.83      210
## 2    10  30      AA    178   JFK  BOS       187 59.00 46.94 64.22      240
## 3    12  18      DL   1585   LGA  MCO       950 33.98 17.06 49.51      270
## 4    11  20      WN   3494   EWR  MDW       711 37.04 17.96 45.58       20
## 5    10  21      DL   2231   LGA  DTW       502 62.96 41.00 44.47      160
## 6    11   7      B6     27   EWR  MCO       937 64.40 55.40 77.29      240
##   wind_speed precip pressure visib                                name      lat
## 1   13.80936      0   1011.0    10   Hartsfield Jackson Atlanta Intl 33.63672
## 2    9.20624      0   1024.9    10 General Edward Lawrence Logan Intl 42.36435
## 3   17.26170      0   1019.8    10                                Orlando Intl 28.42939
## 4    5.75390      0   1035.6    10                                Chicago Midway Intl 41.78597
## 5   13.80936      0   1016.9    10                                Detroit Metro Wayne Co 42.21244
## 6   16.11092      0   1017.5    10                                Orlando Intl 28.42939
##      lon  alt tz dst      tzone year.y      type
## 1 -84.42807 1026 -5  A America/New_York  2001 Fixed wing multi engine
## 2 -71.00518  19 -5  A America/New_York  2002 Fixed wing multi engine
## 3 -81.30899  96 -5  A America/New_York  2002 Fixed wing multi engine
## 4 -87.75242 620 -6  A  America/Chicago  2006 Fixed wing multi engine
## 5 -83.35339 645 -5  A America/New_York  1992 Fixed wing multi engine
## 6 -81.30899  96 -5  A America/New_York  2006 Fixed wing multi engine
##   manufacturer      model engines seats      engine sched_arr_time_minute
## 1          BOEING   737-7H4      2   140 Turbo-fan                25
## 2          BOEING A320-232      2   149 Turbo-fan                 0
## 3          AIRBUS A319-114      2   145 Turbo-fan                20
## 4          BOEING   737-7H4      2   140 Turbo-fan                 5
## 5 AIRBUS INDUSTRIE A320-211      2   182 Turbo-jet               31
## 6          AIRBUS A320-232      2   200 Turbo-fan               57
##   sched_arr_time_num_minute sched_dep_time_minute sched_dep_time_num_minute
## 1                      505                      0                      360
## 2                      840                      50                      770
## 3                     1220                      15                     1035
## 4                     1325                      30                     1230
## 5                     1111                      25                      985
## 6                      717                      0                      540
##   sched_air_time dep_delay_flag temp_flag wind_dir_flag wind_speed_flag
## 1             145              0          0              0              0
## 2              70              0          0              0              0
## 3             185              0          0              0              0
## 4              95              0          0              0              0
## 5             126              0          0              0              0
## 6             177              0          0              0              0
##   precip_flag pressure_flag name_flag year.y_flag type_flag
## 1           0              0          0          0          0
## 2           0              0          0          1          1
## 3           0              0          0          0          0
## 4           0              0          0          0          0
## 5           0              0          0          0          0
## 6           0              1          0          0          0

```

```
library(dplyr)
```

```
DF <- DF %>% mutate_if(is.numeric, scale)
```

```
head(DF)
```

```
##      month      day carrier flight origin dest  distance      temp
## 1 1.30322 -0.9929373    WN   1716    LGA  ATL -0.3777852  0.3339858
## 2 1.01019  1.6325235    AA    178    JFK  BOS -1.1644742  0.1127815
## 3 1.59625  0.2627179    DL   1585    LGA  MCO -0.1205721 -1.2848272
## 4 1.30322  0.4910188    WN   3494    EWR  MDW -0.4475611 -1.1138966
## 5 1.01019  0.6051693    DL   2231    LGA  DTW -0.7335054  0.3339858
## 6 1.30322 -0.9929373    B6     27    EWR  MCO -0.1383581  0.4144237
##      dewp      humid      wind_dir wind_speed      precip      pressure      visib
## 1  0.7418623  0.9315583  0.07717317  0.4871566 -0.1492223 -0.96830167  0.3664282
## 2  0.2753242  0.2375566  0.36735789 -0.3415806 -0.1492223  1.01596006  0.3664282
## 3 -1.2735821 -0.5125364  0.65754261  1.1087096 -0.1492223  0.28792159  0.3664282
## 4 -1.2269283 -0.7129351 -1.76066338 -0.9631336 -0.1492223  2.54341334  0.3664282
## 5 -0.0325909 -0.7695363 -0.40646802  0.4871566 -0.1492223 -0.12606108  0.3664282
## 6  0.7138700  0.9040226  0.36735789  0.9015253 -0.1492223 -0.04040949  0.3664282
##      name      lat      lon      alt
## 1  Hartsfield Jackson Atlanta Intl -0.4207546  0.3298674  0.48364704
## 2  General Edward Lawrence Logan Intl  1.1190951  1.2378347 -0.60619417
## 3      Orlando Intl -1.3395034  0.5408516 -0.52285973
## 4      Chicago Midway Intl  1.0170501  0.1049977  0.04424731
## 5      Detroit Metro Wayne Co  1.0922942  0.4025621  0.07130394
## 6      Orlando Intl -1.3395034  0.5408516 -0.52285973
##      tz dst      tzone      year.y      type
## 1  0.6826595  A America/New_York -0.08500492 Fixed wing multi engine
## 2  0.6826595  A America/New_York  0.08617407 Fixed wing multi engine
## 3  0.6826595  A America/New_York  0.08617407 Fixed wing multi engine
## 4 -0.2514221  A America/Chicago  0.77089000 Fixed wing multi engine
## 5  0.6826595  A America/New_York -1.62561576 Fixed wing multi engine
## 6  0.6826595  A America/New_York  0.77089000 Fixed wing multi engine
##      manufacturer      model      engines      seats      engine
## 1      BOEING  737-7H4  0.05879311  0.02232546 Turbo-fan
## 2      BOEING  A320-232  0.05879311  0.15869100 Turbo-fan
## 3      AIRBUS  A319-114  0.05879311  0.09808410 Turbo-fan
## 4      BOEING  737-7H4  0.05879311  0.02232546 Turbo-fan
## 5  AIRBUS  INDUSTRIE  A320-211  0.05879311  0.65869797 Turbo-jet
## 6      AIRBUS  A320-232  0.05879311  0.93142905 Turbo-fan
##      sched_arr_time_minute sched_arr_time_num_minute sched_dep_time_minute
## 1      -0.2348938      -1.4325947      -1.36042229
## 2      -1.6716145      -0.3129587       1.23408583
## 3      -0.5222379       0.9570761      -0.58206985
## 4      -1.3842703       1.3080068       0.19628258
## 5       0.1099192       0.5927766      -0.06316823
## 6       1.6041087      -0.7240489      -1.36042229
##      sched_dep_time_num_minute sched_air_time dep_delay_flag temp_flag
## 1      -1.6236293      0.14883297       0       0
## 2      -0.1673447     -0.24317221       0       0
## 3       0.7739125      0.35790240       0       0
## 4       1.4665357     -0.11250381       0       0
## 5       0.5963168      0.04952499       0       0
## 6      -0.9842849      0.31608852       0       0
##      wind_dir_flag wind_speed_flag precip_flag pressure_flag name_flag year.y_flag
```

```
## 1      0      0      0      0      0      0
## 2      0      0      0      0      0      1
## 3      0      0      0      0      0      0
## 4      0      0      0      0      0      0
## 5      0      0      0      0      0      0
## 6      0      0      0      1      0      0
##   type_flag
## 1      0
## 2      1
## 3      0
## 4      0
## 5      0
## 6      0
```

```
DF$dep_delay <- dep_delay_vec
```

```
#take out extreme departure delays
```

```
DF<-DF[DF$dep_delay < 30,]
```

```
set.seed(42)
DF$flight <- NULL
train_index <- sample(1:nrow(DF),size=2*nrow(DF)/3,replace=FALSE)
train_df <- DF[train_index,]
test_df <- DF[-train_index,]
```

13 predicting 0

```
rmse = mean((test_df$dep_delay-0)^2) %>% sqrt()
rmse
```

```
## [1] 8.30571
```

14 predicting the mean

```
rmse = mean((test_df$dep_delay-mean(train_df$dep_delay))^2)%>% sqrt()
rmse
```

```
## [1] 8.299767
```

15 predicting the median

```
rmse = mean((test_df$dep_delay-median(train_df$dep_delay))^2)%>% sqrt()
rmse
```

```
## [1] 8.469257
```

16 linear regression with dep

```
model <- lm(dep_delay ~ ., data=train_df)
model_without_dep <- lm(dep_delay ~ .-dest, data=train_df)
anova(model, model_without_dep)

## Analysis of Variance Table
##
## Model 1: dep_delay ~ month + day + carrier + origin + dest + distance +
##      temp + dewp + humid + wind_dir + wind_speed + precip + pressure +
##      visib + name + lat + lon + alt + tz + dst + tzone + year.y +
##      type + manufacturer + model + engines + seats + engine +
##      sched_arr_time_minute + sched_arr_time_num_minute + sched_dep_time_minute +
##      sched_dep_time_num_minute + sched_air_time + dep_delay_flag +
##      temp_flag + wind_dir_flag + wind_speed_flag + precip_flag +
##      pressure_flag + name_flag + year.y_flag + type_flag
## Model 2: dep_delay ~ (month + day + carrier + origin + dest + distance +
##      temp + dewp + humid + wind_dir + wind_speed + precip + pressure +
##      visib + name + lat + lon + alt + tz + dst + tzone + year.y +
##      type + manufacturer + model + engines + seats + engine +
##      sched_arr_time_minute + sched_arr_time_num_minute + sched_dep_time_minute +
##      sched_dep_time_num_minute + sched_air_time + dep_delay_flag +
##      temp_flag + wind_dir_flag + wind_speed_flag + precip_flag +
##      pressure_flag + name_flag + year.y_flag + type_flag) - dest
##   Res.Df    RSS Df Sum of Sq    F    Pr(>F)
## 1 113488 7240975
## 2 113491 7242239 -3    -1263.7 6.602 0.0001863 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

summary <- round(summary(model)$coefficients,6)
sortedddf <- summary[order(summary[,ncol(summary)]),]
head(sortedddf)

##           Estimate Std. Error    t value Pr(>|t|)
## carrierFL    4.615319    0.887510    5.200301      0
## carrierUA    2.440853    0.484369    5.039237      0
## originJFK    1.037568    0.117779    8.809419      0
## wind_speed   0.291278    0.027929   10.429081      0
## precip       0.206814    0.029257    7.068777      0
## pressure    -0.294549    0.027059  -10.885462      0
head(sortedddf)

##           Estimate Std. Error    t value Pr(>|t|)
## carrierFL    4.615319    0.887510    5.200301      0
## carrierUA    2.440853    0.484369    5.039237      0
## originJFK    1.037568    0.117779    8.809419      0
## wind_speed   0.291278    0.027929   10.429081      0
## precip       0.206814    0.029257    7.068777      0
## pressure    -0.294549    0.027059  -10.885462      0
lm_test_df <- test_df

in_test_but_not_train <- setdiff(unique(lm_test_df$model), unique(train_df$model))
```



```

lm_test_df <- lm_test_df[ !lm_test_df$model %in% in_test_but_not_train, ]

in_test_but_not_train <- setdiff(unique(lm_test_df$dest), unique(train_df$dest))
lm_test_df <- lm_test_df[ !lm_test_df$dest %in% in_test_but_not_train, ]

preds = predict(model, newdata=lm_test_df)

## Warning in predict.lm(model, newdata = lm_test_df): prediction from a rank-
## deficient fit may be misleading

rmse = sqrt(mean((lm_test_df$dep_delay - preds)^2))
rmse

## [1] 7.989994

```

17 gbm

```

set.seed(42)
library(gbm)

## Loaded gbm 2.1.5

model <- gbm(dep_delay ~ ., data=train_df,
             n.trees=1000, shrinkage=0.003) # default shrinkage = 0.1


## Distribution not specified, assuming gaussian ...

preds = predict(model, newdata=test_df, n.trees=1000)
rmse = sqrt(mean((test_df$dep_delay - preds)^2))
rmse

## [1] 8.087401

summary(model)

```



```
##                                var      rel.inf
## sched_dep_time_num_minute sched_dep_time_num_minute 40.6986232
## model                        model 31.6301128
## carrier                      carrier 10.3426814
## dest                         dest 8.9487938
## sched_arr_time_num_minute sched_arr_time_num_minute 3.1929834
## origin                      origin 2.1249675
## month                        month 1.8339690
## dewp                         dewp 0.9995901
## precip                      precip 0.2282788
## day                          day 0.0000000
## distance                    distance 0.0000000
## temp                        temp 0.0000000
## humid                       humid 0.0000000
## wind_dir                    wind_dir 0.0000000
## wind_speed                  wind_speed 0.0000000
## pressure                    pressure 0.0000000
## visib                       visib 0.0000000
## name                        name 0.0000000
## lat                         lat 0.0000000
## lon                         lon 0.0000000
## alt                         alt 0.0000000
## tz                          tz 0.0000000
## dst                         dst 0.0000000
## tzone                       tzone 0.0000000
## year.y                      year.y 0.0000000
## type                        type 0.0000000
## manufacturer                manufacturer 0.0000000
## engines                     engines 0.0000000
## seats                       seats 0.0000000
## engine                      engine 0.0000000
## sched_arr_time_minute      sched_arr_time_minute 0.0000000
## sched_dep_time_minute      sched_dep_time_minute 0.0000000
## sched_air_time             sched_air_time 0.0000000
## dep_delay_flag             dep_delay_flag 0.0000000
## temp_flag                  temp_flag 0.0000000
## wind_dir_flag              wind_dir_flag 0.0000000
## wind_speed_flag            wind_speed_flag 0.0000000
## precip_flag                precip_flag 0.0000000
## pressure_flag              pressure_flag 0.0000000
## name_flag                  name_flag 0.0000000
## year.y_flag                year.y_flag 0.0000000
## type_flag                  type_flag 0.0000000
```

Here, you can see the relative influence for each variable for gbm.

For a gbm, the improvement in the splitting criterion (which is mean squared error for regression) for a given variable is calculated at each step. The relative influence for a given variable is the average of these improvements over all the trees where the aforementioned variable is used.

```
model <- gbm(dep_delay ~ ., data=train_df,
             n.trees=1000, shrinkage=0.01) # default shrinkage = 0.1
```

```
## Distribution not specified, assuming gaussian ...
```


```

preds = predict(model, newdata=test_df, n.trees=1000)
rmse = sqrt(mean((test_df$dep_delay - preds)^2))
rmse

```

```
## [1] 7.980812
```

```
summary(model)
```



```

##          var      rel.inf
## sched_dep_time_num_minute sched_dep_time_num_minute 30.83853917
## model          model 25.95730577
## dest          dest 10.17043661
## carrier      carrier 8.62897727
## month        month 8.07615687
## dewp         dewp 4.43848169
## origin      origin 3.05562665
## sched_arr_time_num_minute sched_arr_time_num_minute 2.54883378
## precip      precip 2.42172931
## pressure    pressure 1.52135832
## humid       humid 1.06685951
## dep_delay_flag dep_delay_flag 0.86763349
## pressure_flag pressure_flag 0.21490352
## temp        temp 0.15460862
## day         day 0.03854943
## distance    distance 0.00000000
## wind_dir    wind_dir 0.00000000
## wind_speed  wind_speed 0.00000000
## visib      visib 0.00000000
## name       name 0.00000000
## lat        lat 0.00000000
## lon        lon 0.00000000
## alt        alt 0.00000000
## tz         tz 0.00000000

```

```

## dst                                dst 0.00000000
## tzone                             tzone 0.00000000
## year.y                            year.y 0.00000000
## type                              type 0.00000000
## manufacturer                      manufacturer 0.00000000
## engines                           engines 0.00000000
## seats                             seats 0.00000000
## engine                            engine 0.00000000
## sched_arr_time_minute             sched_arr_time_minute 0.00000000
## sched_dep_time_minute             sched_dep_time_minute 0.00000000
## sched_air_time                    sched_air_time 0.00000000
## temp_flag                         temp_flag 0.00000000
## wind_dir_flag                     wind_dir_flag 0.00000000
## wind_speed_flag                   wind_speed_flag 0.00000000
## precip_flag                       precip_flag 0.00000000
## name_flag                         name_flag 0.00000000
## year.y_flag                       year.y_flag 0.00000000
## type_flag                         type_flag 0.00000000

rmse = sqrt(mean((test_df$dep_delay - preds)^2))
rmse

## [1] 7.980812

set.seed(42)

x <- 2^seq(5,14, by=1) rmse_vec <- numeric(length(x)) count <- 1 for (val in x) { hboost <- gbm( dep_delay
~ ., data = train_df, n.trees = val, distribution = 'gaussian', shrinkage = 0.01 ) preds = predict(hboost, n.trees
= val, newdata = test_df) mse = mean((test_df$dep_delay - preds) ^ 2) rmse <- sqrt(mse) rmse_vec[count]
<- rmse

print(val) print(rmse) count = count + 1 }

plot(x, rmse_vec)

summary(hboost) class(summary(hboost)) summary <- summary(hboost) write.csv(summary,'16384trees_gbm.csv')
gbm_benchmark<-read_csv('shrinkage_0point01_numtrees_32_to_16384_gbm_benchmark.csv')

## Warning: Missing column names filled in: 'X1' [1]

## Parsed with column specification:
## cols(
##   X1 = col_double(),
##   num_trees = col_double(),
##   rmse = col_double()
## )


gbm_benchmark

## # A tibble: 10 x 3
##       X1 num_trees rmse
##   <dbl>   <dbl> <dbl>
## 1     1       32  8.25
## 2     2       64  8.21
## 3     3      128  8.16
## 4     4      256  8.10
## 5     5      512  8.04
## 6     6     1024  7.98

```

##	7	7	2048	7.94
##	8	8	4096	7.92
##	9	9	8192	7.92
##	10	10	16384	7.92

Above values are for gbm with shrinkage of 0.01 Analysis:



Tuning gbm

Here I plotted root mean squared error (rmse) vs the number of trees for shrinkage of 0.01 and all other variables as default for gbm. You can see that after around 5000 trees, increasing the number of trees further gives diminishing returns.

```
library(EZtune) response <- DFdep_delay_eztune_df <- -DF_eztune_df_dep_delay <- NULL eztune_obj <-
eztune(eztune_df, response, method = "gbm", optimizer = "hjn", fast = TRUE, cross = NULL)
```

```
eztune_obj $n [1] 200
```

```
$n.trees [1] 2001
```

```
$interaction.depth [1] 10
```

```
$n.minobsinnode [1] 7
```

```
$shrinkage [1] 0.001
```

```
$mse [1] 72.68835
```

`modelgbm :: gbm(formula = y ., distribution = "gaussian", data = dat, n.trees = results$n.trees, interaction.depth = results$interaction.depth, n.minobsinnode = results$n.minobsinnode, shrinkage = results$shrinkage)` A gradient boosted model with gaussian loss function. 2001 iterations were performed. There were 42 predictors of which 24 had non-zero influence.

```
sqrt(72) [1] 8.485281 sqrt(72.68) [1] 8.525257
```

References

“Gradient Boosting Machines · UC Business Analytics R Programming Guide.” 2019. http://uc-r.github.io/gbm_regression#h2o.